Journal of Statistical Software

MMMMMM YYYY, Volume VV, Issue II. doi: 10.18637/jss.v000.i00

exdqlm: An R Package for Estimation and Analysis of Flexible Dynamic Quantile Linear Models

Antonio Aguirre
University of California Santa Cruz

Raquel Barata
University of California Santa Cruz

Raquel Prado
University of California Santa Cruz

Bruno Sansó
University of California Santa Cruz

Abstract

We present the R package **exdqlm** for Bayesian quantile regression, with primary emphasis on dynamic state-space quantile models for time series. The package is built around flexible dynamic quantile linear models (exDQLMs), which use the extended asymmetric Laplace (exAL) family, a parametric extension of the asymmetric Laplace (AL) commonly used in quantile regression. The software provides exact posterior inference via Markov chain Monte Carlo (MCMC) and fast approximate posterior inference via variational Bayes (VB), supporting posterior uncertainty quantification while also providing a computationally efficient option for longer time series. It also supports static exAL quantile regression with regularized priors, transfer-function models for nonlinear input effects at a given quantile, post hoc posterior-predictive synthesis across separately fitted quantiles, forecasting, and several quantitative and visual diagnostics for model evaluation.

Keywords: dynamic quantile regression, dynamic models, asymmetric Laplace, variational Bayes, coordinate ascent variational inference, quantile regression.

1. Introduction

Quantile regression has become a widely used alternative to mean-based modeling in static settings, but dynamic quantile methods and accompanying software remain comparatively limited. In Bayesian parametric quantile regression, the asymmetric Laplace (AL) likelihood has played a central role because it is linked to the check-loss function used in classical quantile regression (**??**) and admits convenient mixture representations for posterior computation. These computational advantages have made AL-based models attractive, but the AL can be

restrictive as an error model. More flexible error distributions have been studied extensively in the Bayesian nonparametric literature (**???**); however, the resulting methods can be computationally demanding for dynamic models and longer time series. **?** introduce the extended asymmetric Laplace (exAL), a parametric extension of the AL that preserves its computational advantages while relaxing some of its restrictive features. **?** use this family to develop a flexible dynamic quantile linear model (exDQLM), which is the model implemented in the R package **exdqlm**.

The package **exdqlm** provides a unified workflow for Bayesian inference, forecasting, and model assessment for quantile models fitted one quantile at a time. Its core functionality is dynamic state-space quantile modeling, supported by exact posterior simulation via MCMC and a faster variational Bayes (VB) approximation. The package also includes transfer-function exDQLMs for nonlinear input effects at a given quantile, static exAL quantile regression with shrinkage priors, and a synthesis step that combines anchor-specific posterior predictive draws from separately fitted quantile models into a unified posterior predictive distribution. Forecasting and a range of visual and quantitative diagnostics are available throughout. The package is released under the MIT License and is available from the Comprehensive R Archive Network (CRAN) at `https://CRAN.R-project.org/package=exdqlm`.

The current R software ecosystem for quantile regression is broad, but the capabilities most relevant to this paper are distributed across software with different aims. The frequentist package **quantreg** (**?**) remains the broad general-purpose baseline, covering linear, nonlinear, nonparametric, censored, and dynamic quantile regression, together with bootstrap-based and related uncertainty summaries. Its `dynrq()` function supports dynamic linear quantile regression through time-series regression formulas involving lags, differences, trends, seasonal factors, and harmonic terms, while preserving time-series attributes. However, the dynamic structure enters through the design matrix rather than through latent states: once these regressors are constructed, estimation proceeds through the standard **quantreg** fitting routines for a fixed coefficient vector. In that sense, `dynrq()` is a regression-style time-series interface rather than a latent-state state-space model, so it does not provide filtered or smoothed state summaries.

Most other R packages relevant here are static, or else they target a different multi-quantile objective. Among the static Bayesian options, **bayesQR** (**?**) provides Bayesian quantile regression under the AL working likelihood with an adaptive-lasso option, whereas **pqrBayes** (**?**) emphasizes penalized Bayesian quantile regression with spike-and-slab and horseshoe-family shrinkage priors. The broader Bayesian modeling package **brms** (**?**) is also relevant because it supports AL-based quantile regression within a flexible multilevel, nonlinear, and additive regression framework, and can incorporate some autocorrelation structures. However, it is not organized as a dedicated quantile state-space workflow, and its documentation does not present filtered or smoothed state summaries of the type targeted by **exdqlm**. Packages such as **qrjoint** (**?**), **quantregGrowth** (**?**), and **qrcm** (**?**) instead address a different cross-quantile objective through joint estimation, non-crossing constraints, or coefficient modeling across quantiles. Other packages focus on specialized static settings, including additive or semiparametric modeling in **qgam** (**?**), penalized frequentist fitting across quantiles in **rqPen** (**?**), convolution-smoothed linear quantile regression with efficient gradient-based fitting in **conquer** (**?**), and robust or censored-response modeling in **lqr** (**?**).

Table **??** groups a small set of packages by the software distinctions most relevant for situating **exdqlm** within the current quantile-regression ecosystem.

Package	Primary focus	Temporal structure	Fast approximate inference	Cross-quantile treatment
quantreg	Broad frequentist quantile regression	Regression-style dynamic via `dynrq()`	No	Post hoc rearrangement
bayesQR	Static Bayesian quantile regression under the AL working likelihood	Static	No	None
pqrBayes	Static Bayesian penalized quantile regression	Static	No	None
qrjoint	Joint estimation of non-crossing quantiles in regression settings	Static	No	Joint non-crossing estimation
qgam	Additive and semiparametric quantile regression	Static	No	No dedicated joint/non-crossing treatment
exdqlm	Bayesian dynamic/state-space quantile regression; static exAL regression	Latent-state state-space	Yes	Post hoc predictive synthesis

Table 1: Representative R packages spanning broad, static Bayesian, additive, joint multi-quantile, and dynamic/state-space quantile-regression workflows. Temporal structure distinguishes static packages, regression-style dynamic formulations, and latent-state state-space models. Fast approximate inference denotes a package-facing deterministic posterior approximation, such as variational Bayes. Cross-quantile treatment records any package-specific way of linking information across quantile levels; these entries are not methodologically interchangeable. In particular, post hoc predictive synthesis is distinct from joint non-crossing estimation.

Related implementations also exist outside R. In Python, **statsmodels** (**?**) provides linear quantile regression via iterative reweighted least squares. Julia offers **QuantReg.jl** (**?**), and MATLAB includes tools such as `fitrqlinear` and `fitrqnet` (**?**). These implementations are useful reference points for the broader software ecosystem, but they are better viewed as general or static quantile-regression tools than as direct counterparts to a dedicated package for Bayesian dynamic quantile state-space modeling.

Viewed against this software ecosystem, **exdqlm** is not intended to replace broad general-purpose quantile-regression toolkits, additive quantile-regression packages, or software designed specifically for joint multi-quantile estimation. Rather, it addresses a more specific need: a unified Bayesian workflow for dynamic/state-space quantile modeling that combines exact posterior simulation with a fast variational approximation, while also using the more flexible exAL likelihood, supporting static exAL regression with shrinkage priors ranging from ridge to regularized horseshoe variants, and providing transfer-function models, forecasting, diagnostics, and post hoc posterior-predictive synthesis from separately fitted quantile models. General state-space packages such as **dlm** (**?**) and **dynr** (**?**) remain useful for constructing dynamic components, but they do not provide dedicated quantile-regression workflows.

The remainder of the paper is organized as follows. In Section **??**, we describe the modeling framework, including the algorithms used for posterior inference and our methods for including nonlinear relationships. We also discuss static quantile regression as a special case of the framework. In Section **??** we detail the model diagnostics and forecasting method implemented in the package. Then, in Section **??**, we walk through four step-by-step examples

on how to use the package. Finally, in Section **??**, we present a summary of the package capabilities.

2. Extended dynamic quantile linear models

Consider a set of scalar time-evolving responses, y_t, for times $t = 1, \ldots, T$. For each t, an extended dynamic quantile linear model (exDQLM) for inference on a single p_0-th quantile is defined by

$$\text{Observation equation:} \quad y_t = \mathbf{F}_t'\boldsymbol{\theta}_t + \epsilon_t, \qquad \epsilon_t \sim \text{exAL}_{p_0}(0, \sigma, \gamma) \tag{1}$$

$$\text{System equation:} \quad \boldsymbol{\theta}_t = \mathbf{G}_t\boldsymbol{\theta}_{t-1} + \boldsymbol{\omega}_t, \qquad \boldsymbol{\omega}_t \sim \text{N}(\mathbf{0}, \mathbf{W}_t). \tag{2}$$

Here $\mathbf{F}_t$ is the $q \times 1$ regression vector of the covariates corresponding to the $q \times 1$ parameter vector $\boldsymbol{\theta}_t$ at time t, $\mathbf{G}_t$ is the $q \times q$-dimensional evolution matrix defining the structure of the parameter vector evolution in time, and the normal distribution according to which the system error $\boldsymbol{\omega}_t$ evolves is q-variate. The observational errors ϵ_t of the exDQLM follow an exAL distribution for fixed quantile p_0, whose density is denoted as $\text{exAL}_{p_0}(\cdot)$, and it is such that $\int_{-\infty}^{0} \text{exAL}_{p_0}(\epsilon_t|0, \sigma, \gamma)\text{d}\epsilon_t = p_0$. Thus, the observation equation in Equation (**??**) implies that $\mathbf{F}_t'\boldsymbol{\theta}_t$ corresponds to the p_0-quantile of y_t, even if this is not made explicit in the notation above. Table **??** shows the function parameters that define the structure of the dynamic quantile, specifically $\mathbf{F}_t$ and $\mathbf{G}_t$ for $t = 1, \ldots, T$.

The exAL distribution, introduced by **?**, extends the location-scale mixture representation of the AL (**?**). For fixed p_0, the skewness parameter γ is restricted to an interval (L, U), where L and U are the negative and positive roots of $g(\gamma) = 1 - p_0$ and $g(\gamma) = p_0$, respectively, with $g(\gamma) = 2\Phi(-|\gamma|)\exp(\gamma^2/2)$. The package also exports direct distributional utilities for the exAL family. These include `dexal()`, `pexal()`, `qexal()`, and `rexal()`, together with `get_gamma_bounds()` for recovering this admissible interval. This results in a mixture representation of $\text{exAL}_{p_0}(0, \sigma, \gamma)$ as follows:

$$\int\int_{\mathbb{R}^+\times\mathbb{R}^+} \text{N}(y_t|C(p,\gamma)\sigma|\gamma|s_t + A(p)v_t, \sigma B(p)v_t)\text{Exp}(v_t|\sigma)\text{N}^+(s_t|0,1)\text{d}v_t\text{d}s_t. \tag{3}$$

Here, $p = p(p_0, \gamma) = I(\gamma < 0) + \{[p_0 - I(\gamma < 0)]/g(\gamma)\}$ where $g(\gamma) = 2\Phi(-|\gamma|)\exp(\gamma^2/2)$, and $\Phi(\cdot)$ denotes the standard normal CDF. $A(p)$, $B(p)$, $C(p, \gamma)$ are the functions of p and γ: $A(p) = \frac{1-2p}{p(1-p)}$, $B(p) = \frac{2}{p(1-p)}$, $C(p, \gamma) = [I(\gamma > 0) - p]^{-1}$. Lastly, $\text{Exp}(v|\sigma)$ denotes the exponential distribution with mean σ, and $\text{N}^+(s_t|0, 1)$ denotes a normal distribution truncated to the positive reals with mean 0 and variance 1. This mixture representation of the exAL can be exploited to rewrite the exDQLM as the following hierarchical model for $t = 1, \ldots, T$:

$$y_t|\boldsymbol{\theta}_t, \sigma, \gamma, v_t, s_t \sim \text{N}(y_t|\mathbf{F}_t'\boldsymbol{\theta}_t + C(p,\gamma)\sigma|\gamma|s_t + A(p)v_t, \sigma B(p)v_t) \tag{4}$$

$$v_t, s_t|\sigma \sim \text{Exp}(v_t|\sigma)N^+(s_t|0,1) \tag{5}$$

$$\boldsymbol{\theta}_t|\boldsymbol{\theta}_{t-1}, \mathbf{W}_t \sim \text{N}(\mathbf{G}_t\boldsymbol{\theta}_{t-1}, \mathbf{W}_t). \tag{6}$$

This hierarchical form is leveraged for the posterior inference implemented within the **exdqlm** package.

2.1. Prior specification

Input list	`model`				`PriorSigma`		`PriorGamma`		
Symbol	$\mathbf{F}_{1:T}$	$\mathbf{G}_{1:T}$	$\mathbf{m}_0$	$\mathbf{C}_0$	a_σ	b_σ	m_γ	s_γ	ν_γ
List component	`FF`	`GG`	`m0`	`C0`	`a_sig`	`b_sig`	`m_gam`	`s_gam`	`df_gam`
Default	-	-	$\mathbf{0}_q$	$1e3\mathbf{I}_q$	2.1	1.1	0	1	1

Table 2: Translation from mathematical symbols to parameters used by **exdqlm** functions, as well as default values where applicable. $\mathbf{0}_q$ denotes a q-dimensional vector of zeros, and $\mathbf{I}_q$ denotes the q-dimensional identity matrix.

Under a Bayesian formulation and given p_0, we specify priors for the initial state vector $\boldsymbol{\theta}_0$, scale parameter σ and skewness parameter γ. For $\boldsymbol{\theta}_0$, we assume a q-variate conjugate normal prior $\boldsymbol{\theta}_0 \sim \mathrm{N}(\mathbf{m}_0, \mathbf{C}_0)$. For σ, we assume a conjugate inverse-gamma prior denoted as $\mathrm{IG}(a_\sigma, b_\sigma)$. In many cases a strong prior on σ is necessary to guarantee fast posterior convergence and estimation. This is illustrated in the examples of Section **??**. The parameter γ has bounded support over the interval (L, U) where L is the negative root of $g(\gamma) = 1 - p_0$ and U is the positive root of $g(\gamma) = p_0$ (**?**). In contrast to the flat prior suggested by **?**, we find that using a proper prior on the skewness parameter promotes reliable posterior inference. Thus, we place a Student-t distribution truncated to the interval (L, U) as the prior for γ, i.e. $\gamma \sim \mathrm{t}_{(L,U)}(m_\gamma, s_\gamma)$ with ν_γ degrees of freedom. The prior parameters used in the **exdqlm** functions, their corresponding mathematical symbols, and default values are shown in Table **??**.

2.2. Posterior estimation

The hierarchical representation of the exDQLM facilitates posterior simulation using Markov Chain Monte Carlo (MCMC) algorithms. This is implemented in the function `exdqlmMCMC()`. Gibbs sampling steps are used to update the sets of latent variables s_t and v_t for $t = 1, \ldots, T$. The dynamic regression coefficients are sampled using a forward filtering backwards sampling algorithm (**??**). The scale and skewness parameters are updated separately from the latent-state blocks. By default, the package uses a slice-sampling update for the (σ, γ) block, while joint random-walk Metropolis-Hastings (MH) alternatives are also available. When desired, the chain can be warm-started from the variational approximation; this initialization is separate from the subsequent MCMC proposal kernel. An example of this functionality can be found in Section **??**. We perform a total of $N_{BURN} + N_{MCMC}$ MCMC iterations, discarding the first N_{BURN} and retaining the remaining N_{MCMC} posterior samples. The object created by `exdqlmMCMC()` includes posterior samples of all parameters $(\boldsymbol{\theta}_{1:T}, v_{1:T}, s_{1:T}, \gamma, \sigma)$, filtered and smoothed summaries for the state vector $\boldsymbol{\theta}_{1:T}$, and diagnostics for the proposal kernel used in the (σ, γ) update.

For purposes of model selection or large time series data sets, posterior sampling with the MCMC algorithm can be computationally demanding. To address this, the package includes a variational Bayes (VB) routine for fast approximate posterior estimation. The default routine, `exdqlmLDVB()`, uses a mean-field approximation together with a Laplace-Delta treatment of the non-conjugate (σ, γ) block. The variational Bayes (VB) framework approximates the posterior distribution with partitioned variational distributions (**??**). Under a mean-field factorization, this induces a blockwise coordinate-ascent variational inference scheme. For the exDQLM, the blocks associated with the dynamic coefficients and latent variables admit

exdqlm function	`exdqlmMCMC()`				`exdqlmLDVB()`	
Symbol	$\mathcal{K}_{\sigma,\gamma}$	Σ_{MH}	N_{BURN}	N_{MCMC}	ϵ_{tol}	N_{SAMP}
Input	`mh.proposal`	`Sig.mh`	`n.burn`	`n.mcmc`	`tol`	`n.samp`
Default	`"slice"`	$0.05\mathbf{I}_2$	2000	1500	0.1	200

Table 3: Translation from mathematical symbols to parameters used in functions `exdqlmMCMC()` and `exdqlmLDVB()`, as well as default values.

recognizable closed-form updates, while the joint variational distribution of σ and γ is non-conjugate. Following the non-conjugate VB strategy of **?**, we approximate this block on transformed coordinates and use a second-order Delta method to evaluate the expectations of functions of (σ, γ) needed by the remaining variational updates. The dynamic regression coefficients are updated within each iteration using a forward filtering backwards smoothing algorithm. Convergence is assessed jointly through the state, scale, skewness, and ELBO diagnostics. After the algorithm reaches convergence with tolerance ϵ_{tol}, N_{SAMP} samples are drawn from the variational distributions, resulting in an approximate sample of the posterior distribution. The output of `exdqlmLDVB()` includes the approximate posterior samples as well as the parameters of the variational distributions for all parameters $(\boldsymbol{\theta}_{1:T}, v_{1:T}, s_{1:T}, \gamma, \sigma)$. For the (γ, σ) block, the return also contains the transformed mode, approximate covariance matrix, derived moments, and ELBO terms associated with the Laplace-Delta approximation. Selected state-space computations in the dynamic routines are implemented in C++, with corresponding R code retained for parity checks and fallback use.

Table **??** shows the input parameters of the fitting functions `exdqlmMCMC()` and `exdqlmLDVB()`, their mathematical symbols, and their default values.

The argument `Sig.mh` is only used when `mh.proposal` selects a joint random-walk MH kernel; it is not used under the default slice-sampling update.

2.3. Specification of the evolution covariance via discount factors

Both estimation algorithms, `exdqlmMCMC()` and `exdqlmLDVB()`, utilize discount factors to specify the time-evolving covariance matrices $\mathbf{W}_t$ introduced in Equation (**??**) (**?**). More precisely, we define the evolution variance matrix $\mathbf{W}_t = \frac{1-\delta}{\delta}\mathbf{G}_t\mathbf{C}_{t-1}\mathbf{G}_t'$ where $\mathbf{C}_{t-1}$ denotes the posterior variance of the state vector $\boldsymbol{\theta}_{t-1}$ at time $t-1$ and δ is a discount factor such that $0 < \delta \leq 1$. The discount factor δ can be specified by the user via function parameter `df`. Selection of discount factors is typically done by optimizing a model-checking criterion. In practice we recommend predictive criteria such as the CRPS or pplc discussed in Section **??**, with the KL divergence retained as a complementary calibration diagnostic. A brief example of this method for discount factor selection can be found in Section **??**.

Discounting by component of the state-space model is also a functionality of our estimation algorithms. Suppose the state vector is comprised of h components $\boldsymbol{\theta}_{it}$ of dimension q_i for $i = 1, \ldots, h$ such that $\boldsymbol{\theta}_t' = (\boldsymbol{\theta}_{1t}', \ldots, \boldsymbol{\theta}_{ht})'$ and $\sum_{i=1}^{h} q_i = q$. If $\mathbf{W}_{it}$ denotes the ith component of the evolution covariance matrix such that $\mathbf{W}_t = \text{blockdiag}\{\mathbf{W}_{1t}, \ldots, \mathbf{W}_{ht}\}$ and δ_i denotes a discount factor for the ith component, we can define the evolution variance matrix by component as $\mathbf{W}_{it} = \frac{1-\delta_i}{\delta_i}\mathbf{G}_{it}\mathbf{C}_{i,t-1}\mathbf{G}_{it}'$. Here $\mathbf{G}_{it}$ and $\mathbf{C}_{i,t-1}$ denote the ith components of $\mathbf{G}_t$ and $\mathbf{C}_{t-1}$, respectively. The function parameters `df` and `dim.df` allow users to specify a

set of component discount factors $(\delta_1, \ldots, \delta_h)$ and their dimensions $(q_1, \ldots, q_h)$, respectively, within the state-space model. Examples of this feature can be found in Section **??**. For a full review of discount factors including discounting by component, see **?**.

2.4. Transfer function model

Quantifying the relationship between an input and a quantile of a response is a non-trivial task, particularly when the relationship is nonlinear. In mean-centric dynamic modeling, transfer functions are a straightforward method to incorporate variables which measure the combined effect of current and past regression effects (**?**). Within **exdqlm**, this extension is handled by augmenting the dynamic state-space representation conditional on a fixed transfer-function rate λ, and the resulting workflow is available through both VB and MCMC fitting routines.

For $t = 1, \ldots, T$ and transfer input vector $\mathbf{x}_t \in \mathbb{R}^k$, a transfer-function exDQLM with exponential decay is defined by

$$y_t|\boldsymbol{\theta}_t, \zeta_t, \gamma, \sigma \sim \mathrm{exAL}_{p_0}(\mathbf{F}_t'\boldsymbol{\theta}_t + \zeta_t, \sigma, \gamma) \tag{7}$$

$$\boldsymbol{\theta}_t|\boldsymbol{\theta}_{t-1}, \mathbf{W}_t^{\theta} \sim \mathrm{N}(\mathbf{G}_t\boldsymbol{\theta}_{t-1}, \mathbf{W}_t^{\theta}) \tag{8}$$

$$\zeta_t|\zeta_{t-1}, \boldsymbol{\psi}_{t-1}, \omega_t \sim \mathrm{N}(\lambda\zeta_{t-1} + \mathbf{x}_t'\boldsymbol{\psi}_{t-1}, \omega_t) \tag{9}$$

$$\boldsymbol{\psi}_t|\boldsymbol{\psi}_{t-1}, \mathbf{W}_t^{\psi} \sim \mathrm{N}(\boldsymbol{\psi}_{t-1}, \mathbf{W}_t^{\psi}). \tag{10}$$

Here $\zeta_t \in \mathbb{R}$ denotes the accumulated transfer effect on the quantile at time t, while $\boldsymbol{\psi}_t \in \mathbb{R}^k$ is the vector of instantaneous transfer coefficients associated with the components of $\mathbf{x}_t$. The mean contribution of the transfer block is therefore $\lambda\zeta_{t-1} + \mathbf{x}_t'\boldsymbol{\psi}_{t-1}$: the inner product $\mathbf{x}_t'\boldsymbol{\psi}_{t-1}$ captures the contemporaneous effect of the current inputs, and the parameter $\lambda \in [0, 1]$ controls how much of the previous accumulated effect is propagated forward. For $0 < \lambda < 1$, this propagated contribution decays geometrically. More explicitly, the contribution of $\mathbf{x}_t$ to the quantile at time $t + h$ is $\lambda^h\mathbf{x}_t'\boldsymbol{\psi}_{t-1}$. Therefore, for each t and tolerance $\epsilon > 0$, any

$$k_t \geq \frac{\log(\epsilon) - \log(|\mathbf{x}_t'\boldsymbol{\psi}_{t-1}|)}{\log(\lambda)}$$

provides a lower bound on the number of time steps until the magnitude of the propagated effect falls below ϵ. The package reports the median of this bound over time for $\epsilon = 1\mathrm{e}{-3}$.

Conditional on a fixed λ, say $\tilde{\lambda}$, we augment the dynamic fitting routines to incorporate the transfer-function structure by replacing $\mathbf{F}_t$, $\boldsymbol{\theta}_t$, $\mathbf{G}_t$, and $\mathbf{W}_t^{\theta}$ in Equations (**??**)-(**??**) with

$$\tilde{\mathbf{F}}_t' = (\mathbf{F}_t', 1, \mathbf{0}_k'), \qquad \tilde{\boldsymbol{\theta}}_t' = (\boldsymbol{\theta}_t', \zeta_t, \boldsymbol{\psi}_t'),$$

$$\tilde{\mathbf{G}}_t = \mathrm{blockdiag}\{\mathbf{G}_t, \mathbf{G}_t^{tf}\}, \qquad \mathbf{G}_t^{tf} = \begin{pmatrix} \tilde{\lambda} & \mathbf{x}_t' \\ \mathbf{0}_k & \mathbf{I}_k \end{pmatrix},$$

and

$$\tilde{\mathbf{W}}_t = \mathrm{blockdiag}\{\mathbf{W}_t^{\theta}, \omega_t, \mathbf{W}_t^{\psi}\}.$$

This augmentation is implemented in the functions `exdqlmTransferLDVB()` and `exdqlmTransferMCMC()`, which extend the routines found in `exdqlmLDVB()` and `exdqlmMCMC()`.

Symbol	$\mathbf{x}_t$	$\tilde{\lambda}$	$(\delta_\zeta, \boldsymbol{\delta}_\psi)$	$\mathbf{m}_0^{tf}$	$\mathbf{C}_0^{tf}$
transfer-function inputs	`X`	`lam`	`tf.df`	`tf.m0`	`tf.C0`

Table 4: Translation from mathematical symbols to parameters used in the transfer-function fitting routines.

Similar to the specification of $\mathbf{W}_t^\theta$ discussed in Section **??**, discount factors are used to specify the transfer block variances ω_t and $\mathbf{W}_t^\psi$. In the package, `tf.df` can be supplied as a single shared discount factor, as a pair $(\delta_\zeta, \delta_\psi)$ with a shared value for the whole $\boldsymbol{\psi}_t$ block, or componentwise across $(\zeta_t, \psi_{1,t}, \ldots, \psi_{k,t})$. A conjugate normal prior on $(\zeta_0, \boldsymbol{\psi}_0')' \sim N(\mathbf{m}_0^{tf}, \mathbf{C}_0^{tf})$, with $\mathbf{m}_0^{tf} \in \mathbb{R}^{k+1}$ and $\mathbf{C}_0^{tf} \in \mathbb{R}^{(k+1)\times(k+1)}$, completes the transfer-function model extension. Table **??** shows the additional parameters used in the transfer-function fitting routines, as well as their mathematical symbols discussed in this section.

2.5. Special cases, including static quantile regression

The first special case of our exDQLM arises from the fact that the exAL is an extension of the AL (**?**). When $\gamma = 0$, the observational error ϵ_t in Equation (**??**) reduces such that $\epsilon_t \sim \text{exAL}_{p_0}(\epsilon_t|0, \sigma, \gamma) = \text{AL}_{p_0}(\epsilon_t|0, \sigma)$. This special case with observational errors distributed independently from an AL is the dynamic quantile linear model (DQLM) presented in **?**. The DQLM can be implemented with our package using the parameter `dqlm.ind = TRUE`.

Non-time-varying quantile regression is also a special case of our methodology. More specifically, when discount factors are set to 1, the evolution covariance $\mathbf{W}_t$ reduces to $\mathbf{0}$ resulting in non-time-varying parameters (**?**). Thus, the model reduces to the static exAL regression methodology of **?**. The static exAL special case is implemented directly in the package. This model can be implemented using our functions `exalStaticMCMC()` or `exalStaticLDVB()`, which provide MCMC and VB algorithms, respectively.

The final special case is the intersection of the previous two. A static quantile regression model with observational errors distributed according to an AL corresponds to Bayesian linear quantile regression under the AL working likelihood, first presented in **?**. This methodology, utilized to create the Bayesian package for quantile regression **bayesQR**, can also be implemented with our package using functions `exalStaticMCMC()` or `exalStaticLDVB()` with parameter `dqlm.ind = TRUE`, which acts as a convenience flag fixing $\gamma = 0$.

Examples of these special cases can be found in Section 4. A static quantile regression example specifically can be found in Section **??**. For further details on the methodology outlined in this section, see **?**.

3. Model Diagnostics and Forecasting

To evaluate the quantile inference and predictive performance of the exDQLM, several quantitative and visual model diagnostics are included in the package.

The one-step-ahead predictive distribution function introduced by **?** is a useful diagnostic for models in which the marginal forecast distributions are unrecognizable or difficult to compute

(**??**). For the exDQLM, this distribution is defined as

$$u_t = \Pr(Y_t \leq y_t | y_{1:t-1}, v_{1:T}, s_{1:T}, \sigma, \gamma), \tag{11}$$

with $y_{1:t-1} = y_1, \ldots, y_{t-1}$, $v_{1:T} = (v_1, \ldots, v_T)$ and $s_{1:T} = (s_1, \ldots, s_T)$. Here u_t defines an independent sequence which is uniformly distributed on the interval $(0, 1)$ (**?**). Conditional on $\{v_{1:T}, s_{1:T}, \sigma, \gamma\}$, the forecast distribution of y_t within the exDQLM is normally distributed. Therefore, conditional on MAP estimates $\{\hat{v}_{1:T}, \hat{s}_{1:T}, \hat{\sigma}, \hat{\gamma}\}$ we estimate the sequence as

$$\hat{u}_t = \Phi(y_t | y_{1:t-1}, \hat{v}_{1:T}, \hat{s}_{1:T}, \hat{\sigma}, \hat{\gamma}) \tag{12}$$

where Φ denotes the normal CDF.

Using this estimated sequence $\{\hat{u}_t\}$, we can assess the model performance in several ways. First, we visually examine the correlation of the estimated sequence via an ACF plot. Second, by transforming the values with a standard normal inverse CDF, we compare their distributional shape to that of a standard normal with a normal QQ-plot. To quantify the divergence of this transformed sequence from normality we compute the KL divergence (**?**), $\mathrm{KL}(h, \phi) = \int_{-\infty}^{\infty} h(x) \log \frac{h(x)}{\phi(x)} \mathrm{d}x$, where h denotes the numerically approximated density of our transformed sample and ϕ is the standard normal density. A smaller KL divergence suggests a superior model fit. Lastly, the MAP standardized forecast errors,

$$\frac{y_t - \mathsf{E}[y_t | y_{1:t-1}, \hat{v}_{1:T}, \hat{s}_{1:T}, \hat{\sigma}, \hat{\gamma}]}{\sqrt{\mathrm{Var}(y_t | y_{1:t-1}, \hat{v}_{1:T}, \hat{s}_{1:T}, \hat{\sigma}, \hat{\gamma})}} \tag{13}$$

can also be used as an additional graphical aid in examining the distribution at specific time points. These MAP standardized forecast errors are included in the output of the `exdqlmMCMC()` and `exdqlmLDVB()` functions.

To connect the predictive criteria used below, let $\rho_p(x) = x[p - I(x < 0)]$ denote the check loss at level p. In addition to the calibration-based diagnostics above, we compute the mean continuous ranked probability score (CRPS) from the posterior predictive draws (**?**). For a predictive distribution G with quantile function G^{-1} and an observation y, the CRPS is

$$\mathrm{CRPS}(y, G) = \mathbb{E}_{Y \sim G} |Y - y| - \frac{1}{2} \mathbb{E}_{Y, Y' \sim G} |Y - Y'| = 2 \int_0^1 \rho_p\Big(y - G^{-1}(p)\Big) \, dp,$$

This representation shows that the CRPS aggregates check-loss performance over all quantile levels, so it evaluates the full predictive distribution rather than a single target quantile. It rewards forecasts that are both sharp and well calibrated, and it reduces to the absolute error, and hence to the MAE after averaging, for deterministic forecasts. Smaller values indicate better predictive performance.

For a final diagnostic, we include **?** posterior predictive loss criterion (pplc), which evaluates the posterior predictive distribution under the target-level check loss ρ_{p_0}. That is,

$$\mathrm{pplc} = \sum_t \mathsf{E}[\rho_{p_0}(y_t^{obs} - y_t^{rep}) | y_{1:T}], \tag{14}$$

where the expectation is with respect to the posterior predictive distribution of y_t, $p(y_t^{rep} | y_{1:T})$. This distribution is used to generate a replicate dataset and is the expected value of the observation equation (**??**) with respect to the posterior distributions of the parameters. Thus, the

Symbol	$\tilde{t}$	k	$\mathbf{F}_{(\tilde{t}+1):(\tilde{t}+k)}$	$\mathbf{G}_{(\tilde{t}+1):(\tilde{t}+k)}$
`exdqlmForecast()` input	`start.t`	`k`	`fFF`	`fGG`

Table 5: Translation from mathematical symbols to additional parameters used in function `exdqlmForecast()`.

CRPS summarizes the full predictive distribution through an integrated check loss, whereas the pplc keeps the target quantile explicit through ρ_{p_0}. Smaller pplc values suggest a superior model fit. Samples from the posterior replicate distributions are included in the output of the `exdqlmMCMC()` and `exdqlmLDVB()` functions.

The function `exdqlmDiagnostics()` implements the model checking criteria discussed in this section. When an output from either `exdqlmMCMC()` and `exdqlmLDVB()` is passed to `exdqlmDiagnostics()`, the resulting object includes: the estimated sequence $\{\hat{u}_t\}$, the KL divergence, the mean CRPS, pplc, the ordered pairs of the QQ-plot, and autocorrelations by lag. By default, the function also produces the QQ-plot, ACF plot and MAP standard forecast error plot. Examples of the utility of `exdqlmDiagnostics()` can be found in the following Section **??**.

3.1. k-step-ahead forecast

For an arbitrary time $\tilde{t}$, the k-step-ahead future marginal distribution of the quantile is

$$\mathbf{F}'_{\tilde{t}+k}\boldsymbol{\theta}_{\tilde{t}+k}|y_1,\ldots,y_{\tilde{t}} \sim \mathrm{N}(\mathbf{F}'_{\tilde{t}+k}a_{\tilde{t}}(k), \mathbf{F}'_{\tilde{t}+k}\mathbf{R}_{\tilde{t}}(k)\mathbf{F}_{\tilde{t}+k}) \tag{15}$$

where $\mathbf{a}_{\tilde{t}}(k) = \mathbf{G}_{\tilde{t}+k}\mathbf{a}_{\tilde{t}}(k-1)$, $\mathbf{R}_{\tilde{t}}(k) = \mathbf{G}_{\tilde{t}+k}\mathbf{R}_{\tilde{t}}(k-1)\mathbf{G}'_{\tilde{t}+k} + \mathbf{W}_{\tilde{t}+k}$, $\mathbf{a}_{\tilde{t}}(0) = \mathbf{m}_{\tilde{t}}$, and $\mathbf{R}_{\tilde{t}}(0) = \mathbf{C}_{\tilde{t}}$, with $\mathbf{m}_{\tilde{t}}$ and $\mathbf{C}_{\tilde{t}}$ denoting the filtered mean and covariance of $\boldsymbol{\theta}_{\tilde{t}}$, respectively. This distribution is implemented in the function `exdqlmForecast()`, which is compatible with the outputs from both the MCMC and VB algorithms. Table **??** shows the input parameters used in the function, as well as their mathematical symbols seen in this section. The output of `exdqlmForecast()` includes the parameters of the forecasted distribution as well as $\mathbf{a}_{\tilde{t}}(k)$ and $\mathbf{R}_{\tilde{t}}(k)$ for the k steps.

The manuscript figures and tables are generated from scripts distributed with the article repository. These scripts rebuild the example artifacts directly from the current **exdqlm** API, and the accompanying reproducibility materials record the computing environment and manuscript outputs used in the paper.

4. Examples

Table **??** provides an overview of the **exdqlm** functions and the resulting objects. We demonstrate the general workflow and key features of the package through several analyses on real datasets. The first three examples illustrate dynamic workflows, while the fourth illustrates the static exAL special case. The following examples are implemented on a personal computer with an Apple M1 processor. We begin by loading the **exdqlm** package used for all examples.

```
R> library(exdqlm)
```

Function	Output	Description
Model construction		
`polytrendMod()`	`exdqlm`	Create n-th order polynomial component.
`seasMod()`	`exdqlm`	Create Fourier form periodic component.
`regMod()`	`exdqlm`	Create standard regression component.
`as.exdqlm()`	`exdqlm`	Coerce a `list` into an `exdqlm` object.
Model fitting – Input `exdqlm` objects		
`exdqlmMCMC()`	`exdqlmMCMC`	Estimate the posterior with MCMC algorithm.
`exdqlmLDVB()`	`exdqlmLDVB`	Estimate the posterior with a VB approximation.
`exdqlmTransferLDVB()`	`exdqlmLDVB`	Estimate the transfer-function exDQLM with a VB approximation.
`exdqlmTransferMCMC()`	`exdqlmMCMC`	Estimate the transfer-function exDQLM with MCMC.
Static fitting – Input design matrix and response		
`exalStaticLDVB()`	`exalStaticLDVB`	Estimate static exAL regression with a VB approximation.
`exalStaticMCMC()`	`exalStaticMCMC`	Estimate static exAL regression with MCMC.
Model fit exploration – Input `exdqlmMCMC` or `exdqlmLDVB` objects		
`exdqlmForecast()`	`exdqlmForecast`	Compute k-step-ahead forecast quantiles.
`exdqlmDiagnostics()`	`exdqlmDiagnostic`	Compute & plot model diagnostics.
`exdqlmPlot()`	`list`	Plot posterior summary of dynamic quantile.
`compPlot()`	`list`	Plot posterior summary of component/parameter.
`exalStaticDiagnostics()`	`exalStaticDiagnostic`	Compute static fit diagnostics and comparisons.
Predictive synthesis – Input posterior predictive draw lists		
`quantileSynthesis()`	`list`	Construct a unified posterior predictive distribution from anchor-specific predictive draws across fitted quantiles.

Table 6: Overview of the core workflow functions used in this paper. Generic S3 methods defined for the fitted objects listed in the table include `plot()`, `print()`, and `summary()`.

4.1. Lake Huron

In this example, we consider the Lake Huron time series from the package **datasets**. The data are annual measurements of the level (ft) of Lake Huron from 1875 to 1972, shown in Figure **??**. With this example we show how to: build a basic state-space structure, specify a prior on γ, run the MCMC algorithm, plot the estimated quantiles and forecast distributions.

To estimate the dynamic distribution of the data, we consider three quantiles, $p_0 = 0.95, 0.50,$ and 0.05. We begin by creating the state-space structure and prior used to estimate the quantiles (i.e. $\mathbf{F}_t, \mathbf{G}_t, \mathbf{m}_0$, and $\mathbf{C}_0$ discussed in Section **??**). We choose to model the quantiles with a second order polynomial trend, which we construct with the `polytrendMod()` function.

```
R> model = polytrendMod(order = 2, m0 = c(mean(LakeHuron),0), C0 = 10*diag(2))
R> model

Prior mean of the state vector:
$m0
         [,1]
[1,] 579.0041
[2,]   0.0000


Prior covariance of the state vector:
$C0
     [,1] [,2]
[1,]   10    0
[2,]    0   10


Observational vector:
$FF
     [,1]
[1,]    1
[2,]    0


Evolution matrix:
$GG
     [,1] [,2]
[1,]    1    1
[2,]    0    1
```

To allow variability in the dynamic quantile, we select a single discount factor of 0.9 to define the evolution covariance matrix $\mathbf{W}_t$. Discount factors can also be optimized, an example of which we will show in Section **??**. We place truncated Student-t priors on the skewness parameters via `PriorGamma` as discussed in Section **??**. In the fits shown here we use `n.burn = 2000` and `n.mcmc = 3000`. The R output is shown for only the last quantile using the input `verbose = TRUE`.

```
R> M95 = exdqlmMCMC(y = LakeHuron, p0 = 0.95, model = model,
+                 df = 0.9, dim.df = 2,
+                 PriorGamma = list(m_gam = -1, s_gam = 0.1, df_gam = 1),
+                 n.burn = 2000, n.mcmc = 3000, verbose = FALSE)
R> M50 = exdqlmMCMC(y = LakeHuron, p0 = 0.50, model = model,
+                  df = 0.9, dim.df = 2,
+                  PriorGamma = list(m_gam = 0, s_gam = 0.1, df_gam = 1),
+                  n.burn = 2000, n.mcmc = 3000, verbose = FALSE)
R> M5 = exdqlmMCMC(y = LakeHuron, p0 = 0.05, model = model,
+                 df = 0.9, dim.df = 2,
+                 PriorGamma = list(m_gam = 1, s_gam = 0.1, df_gam = 1),
+                 n.burn = 2000, n.mcmc = 3000, verbose = TRUE)
```

Representative console excerpt, abridged to the fields most relevant for the progress format:

```
MCMC backend: C++ (mode=fast)
MCMC init: running LDVB warm start
MCMC start | burn=2000 | keep=3000 | kernel=slice | warm_start=ldvb
```

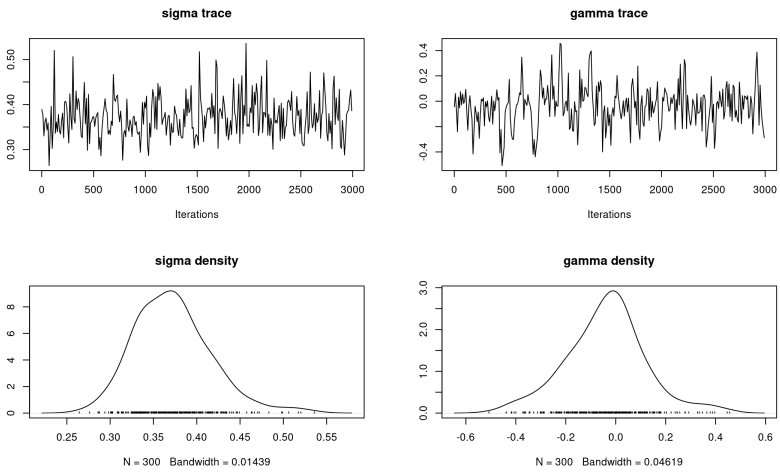

Figure 1: Example 1: Lake Huron. Trace and density plots from the dedicated free-σ median run with 2000 burn-in iterations and 3000 retained draws. For readability, every 10th retained draw is used in the plotting step. The top row shows the posterior trace plots for σ and γ, and the bottom row shows the corresponding posterior densities.

```
...
MCMC done | runtime_sec=122.179 | sigma=0.220 | gamma=5.013 | kernel=slice
```

For this time series of length 98, the free-σ run shown above takes about 122.2 seconds to sample 5000 iterations. By default `exdqlmMCMC()` uses the package's VB routine, `exdqlmLDVB()`, to warm-start the chain, and this initialization choice can be changed independently of the subsequent MCMC kernel. For the free-σ Lake Huron median fit, the default `slice` update was more stable than `laplace_rw` in our support comparison, so we retain `slice` here. Because the data are relatively symmetric, we do not expect the skewness parameter γ to be far from zero at $p_0 = 0.5$; in the dedicated median run with free σ, the posterior mean of σ is about 0.375 with a 95% credible interval $(0.298, 0.473)$, while the posterior mean of γ is -0.044 with a 95% credible interval $(-0.375, 0.321)$. All retained samples in the output of `exdqlmMCMC()` are `mcmc` objects, so we can use the package **coda** to examine the trace and density plots of both σ and γ, seen in Figure **??**.

```
R> M50.trace = exdqlmMCMC(y = LakeHuron, p0 = 0.50, model = model,
+                   df = 0.9, dim.df = 2,
+                   PriorGamma = list(m_gam = 0, s_gam = 0.1, df_gam = 1),
+                   n.burn = 2000, n.mcmc = 3000, verbose = FALSE)
R> keep.idx = seq(1, length(M50.trace$samp.sigma), by = 10)
R> sigma.trace = coda::mcmc(M50.trace$samp.sigma[keep.idx], thin = 10)
R> gamma.trace = coda::mcmc(M50.trace$samp.gamma[keep.idx], thin = 10)
R> coda::traceplot(sigma.trace, main = "sigma trace")
R> coda::densplot(sigma.trace, main = "sigma density")
R> coda::traceplot(gamma.trace, main = "gamma trace")
R> coda::densplot(gamma.trace, main = "gamma density")
```

The posterior summary indicates that γ is not distinguishable from zero, so the additional flexibility of the skewness parameter is not needed for the median fit. We therefore re-run

the model with a point-mass prior on γ at 0 using the settings `gam.init = 0` and `fix.gamma = TRUE`, which is equivalent to the setting `dqlm.ind = TRUE`. This reduces the model to the DQLM mentioned in Section ??.

```
R> M50 = exdqlmMCMC(y = LakeHuron, p0 = 0.50, model = model,
+                   df = 0.9, dim.df = 2,
+                   gam.init = 0, fix.gamma = TRUE,
+                   n.burn = 2000, n.mcmc = 3000, verbose = TRUE)
```

```
MCMC backend: C++ (mode=fast)
MCMC init: running LDVB warm start
MCMC start | burn=2000 | keep=3000 | kernel=conjugate | warm_start=ldvb
...
MCMC done | runtime_sec=51.838 | sigma=0.430 | kernel=conjugate
```

The results are `exdqlmMCMC` objects that can be used for analysis and forecasting. First we examine the estimated quantiles with the plot method for `exdqlmMCMC` objects, seen in the top-left panel of Figure ??. This generates a plot of the data with the MAP estimates and the 95% CrIs of the dynamic quantile. Subsequent quantiles (i.e. $p_0 = 0.50, 0.05$ in this example) can be added to the plot using the function `exdqlmPlot()` with the argument `add = TRUE`.

```
R> plot(M95)
R> exdqlmPlot(M50, add = TRUE, col = "blue")
R> exdqlmPlot(M5, add = TRUE, col = "forest green")
R> legend("topright", lty = 1, col = c("purple","blue","forest green"),
+        legend = c(expression('p'[0]*'=0.95'),expression('p'[0]*'=0.50'),
+                   expression('p'[0]*'=0.05')))
```

Next, we forecast the $k = 8$ step-ahead distributions starting in 1972. To do this, we first create the observational vector (`fFF`) and evolution matrix (`fGG`) that will be used in the forecast updates. In this case both are non-time-varying and identical to those used to estimate the quantile.

```
R> fFF = model$FF
R> fGG =  model$GG
```

We plot the time series data in a narrower range for closer examination and add the forecast estimates using `exdqlmForecast()` with argument `add = TRUE`. Notice we start the forecast at the last time index of the data, i.e. `start.t = length(LakeHuron)`.

```
R> plot(LakeHuron,  xlim = c(1952,1980), ylim = c(575,581),
+        col = "dark grey")
R> exdqlmForecast(start.t = length(LakeHuron), k = 8, m1 = M95,
+               fFF = fFF, fGG = fGG, plot = TRUE, add = TRUE)
R> exdqlmForecast(start.t = length(LakeHuron), k = 8, m1 = M50,
+               fFF = fFF, fGG = fGG, plot = TRUE, add = TRUE,
+               cols = c("blue","light blue"))
R> exdqlmForecast(start.t = length(LakeHuron), k = 8, m1 = M5,
+               fFF = fFF, fGG = fGG, plot = TRUE, add = TRUE,
+               cols = c("forest green","green"))
```

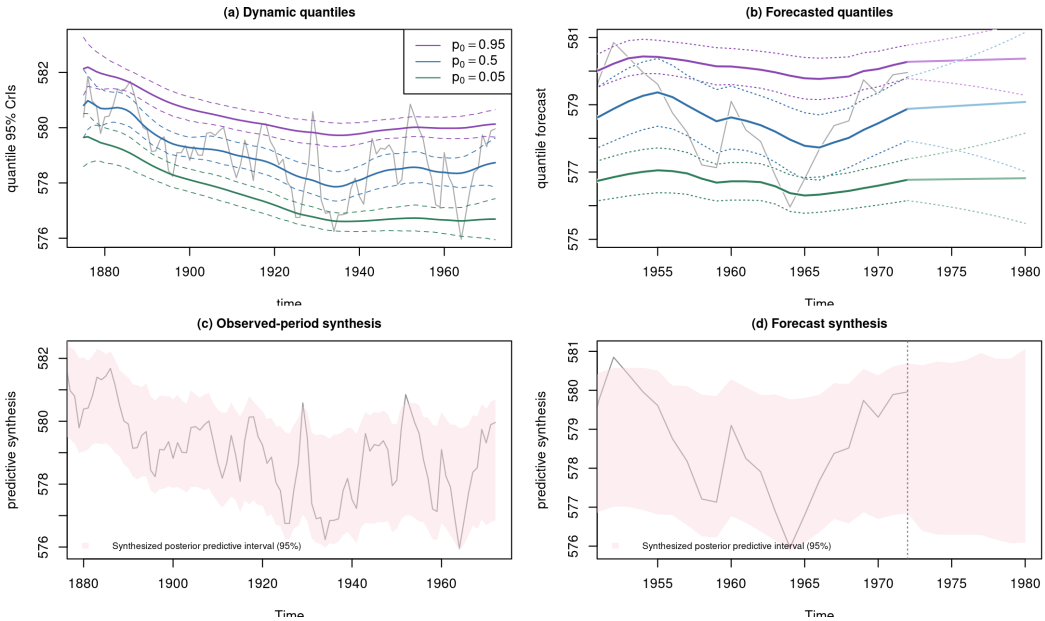

Figure 2: Example 1: Lake Huron. Top-left: MAP estimates and 95% CrIs of the estimated quantiles, plotted with the data (grey). Top-right: forecasted quantile estimates and 95% credible intervals (seen after 1972), along with the filtered quantile estimates and 95% credible intervals (seen from 1952 to 1972). Bottom-left: synthesized posterior predictive 95% interval over the observed period. Bottom-right: synthesized posterior predictive 95% interval over the forecast window.

This generates a plot of the data with the forecasted quantile estimates and 95% credible intervals (seen after 1972), along with the filtered quantile estimates and 95% credible intervals (before 1972), for reference. Results can be seen in the top-right panel of Figure **??**. The percentage in the CrIs can be modified with the parameter `cr.percent`.

When several quantile models are fitted separately, the package also allows their posterior predictive draws to be combined into a single coherent predictive distribution through `quantileSynthesis()`. Each fitted model contributes local predictive information near its target quantile, and the synthesis step interpolates across these anchor-specific draws to obtain one posterior predictive distribution. When independently fitted quantiles induce crossing, optional monotonicity corrections based on isotonic adjustment and monotone rearrangement are available (**??**). For the Lake Huron example, we apply this routine to the posterior predictive samples stored in `M5`, `M50`, and `M95` over the observed period. Since these posterior predictive draws are produced from the smoothed state draws, the lower panels of Figure **??** display the synthesized posterior predictive 95% interval over the observed and forecast windows. For the eight-step forecast horizon, we first generate future predictive draws from each fitted model using the corresponding forecasted quantile distributions together with posterior draws of σ and γ, and then apply the same synthesis step. This provides a single predictive envelope while preserving the separate quantile-specific fits.

A direct call for the observed posterior predictive draws is given below:

```
R> syn.obs = quantileSynthesis(
```

```
+                draws_list = list(M5$samp.post.pred, M50$samp.post.pred,
+                                  M95$samp.post.pred),
+                p = c(0.05, 0.50, 0.95),
+                T_expected = length(LakeHuron))
```

The function `exdqlmForecast()` returns the forecast summaries, but not matrices of future observations. To construct those, we first sample future quantile values from the forecast distribution and then map them through `rexal()` using matched posterior draws of σ and γ. For example, for the $p_0 = 0.05$ fit:

```
R> q.fore = sweep(matrix(rnorm(8 * n.samp), 8, n.samp), 1,
+                 sqrt(fc05$fQ), "*") + fc05$ff
R> future.M5$ydraw = vapply(1:n.samp, function(j)
+                 rexal(8, p0 = M5$p0, mu = q.fore[, j],
+                       sigma = sigma.draws[j], gamma = gamma.draws[j]),
+                 numeric(8))
```

Here `sigma.draws` and `gamma.draws` denote matched retained posterior draws from the fitted model. Analogous matrices are generated for `M50` and `M95`.

We then collect these matrices into `future.draws` and synthesize them in the same way:

```
R> future.draws = list(future.M5$ydraw, future.M50$ydraw,
+                      future.M95$ydraw)

R> syn.fore = quantileSynthesis(
+                draws_list = future.draws,
+                p = c(0.05, 0.50, 0.95),
+                T_expected = 8)
```

4.2. Sunspots

For our next example, we use the yearly Sunspot time series from the package **datasets**. The data are yearly counts for sunspots from 1700 to 1988, shown in Figure **??**. In this example we show how to use the **dlm** package to create the state-space model, combine blocks of a state-space model, apply the VB approximation through `exdqlmLDVB()`, perform visual model diagnostics to compare the exDQLM with the DQLM, and use those diagnostics for discount-factor selection.

The solar cycle impacts astronauts in space as well as lives and technology on earth. Of particular interest to scientists is the solar max, or period in which the sun is most active **?**, therefore we will explore the 0.85 quantile in this example. Again, we begin by building the state-space structure used to estimate the quantile. Here we choose to model the quantile with a first order polynomial trend component as well as a Fourier form seasonal component that incorporates a fundamental period every 11 observations (years), as well as some of its corresponding harmonics. This results in a dynamic model with a total of 9 state parameters, 1 for the first order polynomial component and 8 for the seasonal component. The **exdqlm** functions are compatible with time-invariant `dlm` objects from the **dlm** package. For example, we create the first order trend component with the function `dlmModPoly()` from the **dlm** package, then use `as.exdqlm()` to turn the output into an `exdqlm` object.

```
R> dlm.trend.comp = dlm::dlmModPoly(1, m0 = mean(sunspot.year), C0 = 10)
R> trend.comp = as.exdqlm(dlm.trend.comp)
```

Next, we construct the seasonal component with our `seasMod` function. The duration of sunspot cycles typically follow a period of 11 years, of which we include the first 4 harmonics in the component.

```
R> seas.comp = seasMod(p = 11, h = 1:4, C0 = 10*diag(8))
```

To combine the two components into a single state-space structure, we add the two `exdqlm` objects.

```
R> model = trend.comp + seas.comp
```

The resulting evolution matrix of the combined models is printed for illustration.

```
R> model$GG
```

```
     [,1]    [,2]   [,3]    [,4]   [,5]    [,6]    [,7]    [,8]    [,9]
 [1,]    1  0.0000 0.0000  0.0000 0.0000  0.0000  0.0000  0.0000  0.0000
 [2,]    0  0.8413 0.5406  0.0000 0.0000  0.0000  0.0000  0.0000  0.0000
 [3,]    0 -0.5406 0.8413  0.0000 0.0000  0.0000  0.0000  0.0000  0.0000
 [4,]    0  0.0000 0.0000  0.4154 0.9096  0.0000  0.0000  0.0000  0.0000
 [5,]    0  0.0000 0.0000 -0.9096 0.4154  0.0000  0.0000  0.0000  0.0000
 [6,]    0  0.0000 0.0000  0.0000 0.0000 -0.1423  0.9898  0.0000  0.0000
 [7,]    0  0.0000 0.0000  0.0000 0.0000 -0.9898 -0.1423  0.0000  0.0000
 [8,]    0  0.0000 0.0000  0.0000 0.0000  0.0000  0.0000 -0.6549  0.7557
 [9,]    0  0.0000 0.0000  0.0000 0.0000  0.0000  0.0000 -0.7557 -0.6549
```

We will apply discount factors by component, i.e. a discount factor of 0.9 for the trend component of dimension 1 and a discount factor of 0.85 for the seasonal component of dimension 8. This discounting strategy can be applied with arguments `df = c(0.9,0.85)` and `dim.df = c(1,8)`.

We call the function `exdqlmLDVB()` to obtain the VB approximations. We use the routine with argument `dqlm.ind = TRUE` to estimate the DQLM, as well as with the default settings to estimate the exDQLM. In this example we set `n.samp = 3000`. The output is suppressed using input `verbose = FALSE`, however we check the run-time of each model after convergence using the `summary()` generic function. For this dataset of length 289, the DQLM VB fit completes in about 10 seconds. The corresponding exDQLM VB fit is more computationally demanding and takes about 114 seconds.

```
R> M1 = exdqlmLDVB(y = sunspot.year, p0 = 0.85, model = model,
+                  df = c(0.9,0.85), dim.df = c(1,8),
+                  dqlm.ind = TRUE, fix.sigma = FALSE,
+                  n.samp = 3000,
+                  verbose = FALSE)
R> M2 = exdqlmLDVB(y = sunspot.year, p0 = 0.85, model = model,
+                  df = c(0.9,0.85), dim.df = c(1,8),
+                  fix.sigma = FALSE,
+                  n.samp = 3000,
+                  verbose = FALSE)
R> summary(M1)
```

```
Bayesian Dynamic Quantile Regression Model (DQLM)
Number of Observations: 289
```

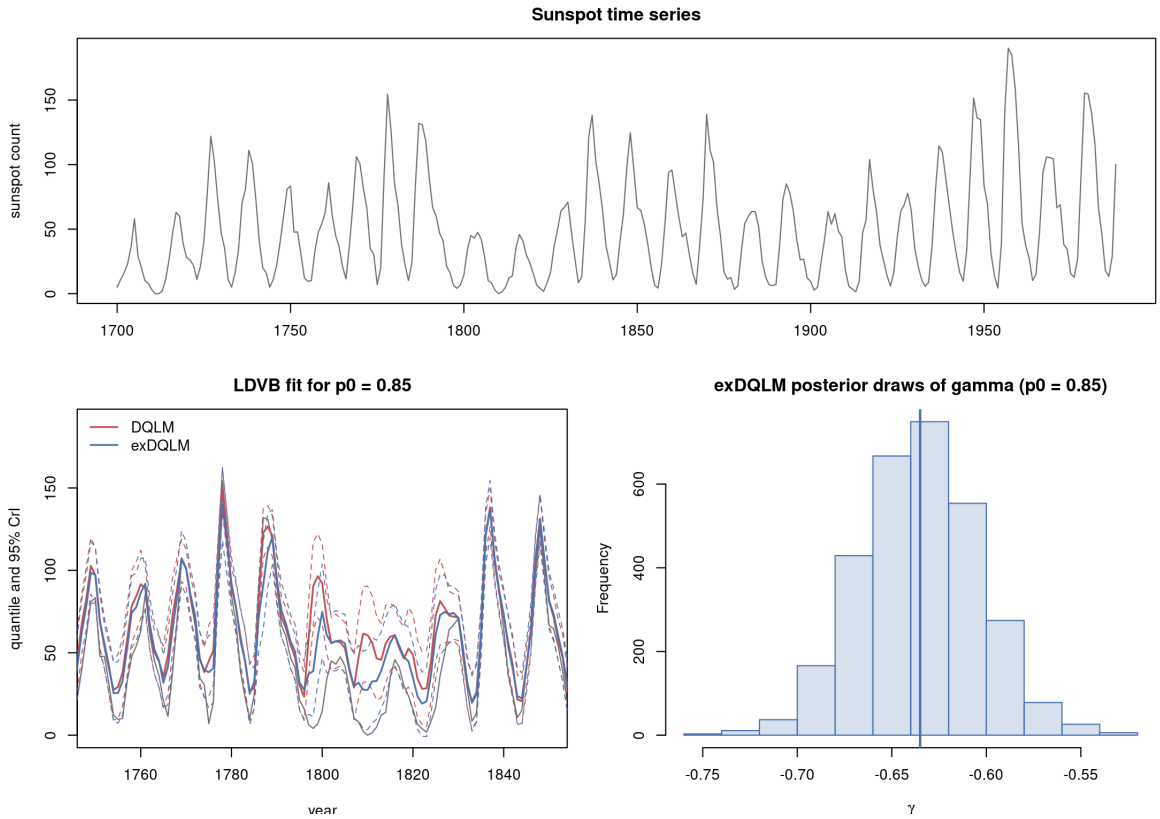

Figure 3: Example 2: Sunspots. Top: The sunspot time series from 1700 to 1988. Bottom left: MAP estimates and 95% CrIs of the estimated quantiles from the DQLM (red) and exDQLM (blue), plotted with the data (grey) from 1750 to 1850. Bottom right: Histogram of samples from the approximated posterior distribution of γ under the exDQLM.

```
State Dimension: 9
Discount factors ( dimensions ): 0.9 ( 1 ), 0.85 ( 8 )

DQLM fitted using VB
Iterations until convergence: 67
Run-time: 9.522 seconds

R> summary(M2)

Bayesian Dynamic Quantile Regression Model (exDQLM)
Number of Observations: 289
State Dimension: 9
Discount factors ( dimensions ): 0.9 ( 1 ), 0.85 ( 8 )

exDQLM fitted using LDVB
Parameters estimated with LD: gamma sigma
Iterations until convergence: 200
Run-time: 113.551 seconds
```

We can visualize the differences between the two resulting dynamic quantiles using the function `exdqlmPlot()`, shown in the lower-left panel of Figure **??**. For clarity we limit that panel

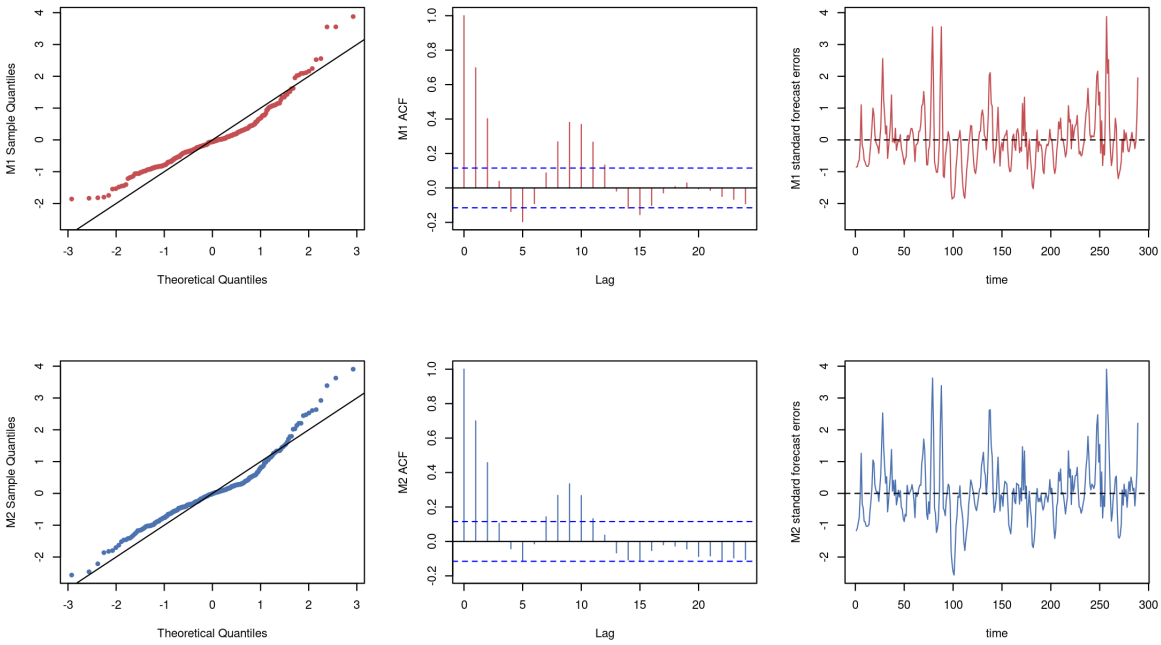

Figure 4: Example 2: Sunspots. The QQ-plots (left column), ACF plots (center column) and standard forecast errors (right column) from the DQLM (red) and exDQLM (blue).

to years 1750 to 1850.

```
R> plot(sunspot.year, xlim = c(1750,1850), col = "dark grey",
+       ylab = "quantile 95% CrIs")
R> exdqlmPlot(M1, add = TRUE, col = "red")
R> exdqlmPlot(M2, add = TRUE, col = "blue")
```

To examine whether the added flexibility of the exDQLM is significant, we can also visualize the approximate posterior distribution of the skewness parameter γ by plotting the samples from the corresponding variational distribution, also seen in Figure **??**.

```
R> hist(M2$samp.gamma,xlab=expression(gamma),main="")
```

The approximate distribution is clearly distinct from zero, thus we expect a better model fit from the exDQLM.

To further examine the two models, we use the function `exdqlmDiagnostics()` to plot the model diagnostics discussed in Section **??**. The results are seen in Figure **??**. The QQ-plot and the ACF of the standard forecast errors both favor the exDQLM.

```
R> exdqlmDiagnostics(M1, M2, cols = c("red","blue"))
```

Finally, the function `exdqlmDiagnostics()` can also be used for model selection. Say we want to check whether the discount factor we used for the seasonal component is reasonable under a predictive scoring rule. We can create a list of possible discount factors, quickly run the VB approximation for each possible model, and compare the resulting CRPS and KL

model	method	runtime (s)	KL	CRPS	pplc
DQLM	VB	**9.52**	0.144	8.588	3909.4
DQLM	MCMC	266.16	**0.136**	**7.664**	**3048.1**
exDQLM	VB	**113.55**	**0.078**	6.424	2346.5
exDQLM	MCMC	330.91	0.111	**5.069**	**2186.7**

Table 7: Example 2: representative dynamic benchmark for the DQLM and exDQLM fits at $p_0 = 0.85$. Each row reports runtime together with the KL divergence, the mean CRPS from posterior predictive draws, and the posterior predictive loss criterion. The VB rows use `n.samp = 3000`. The MCMC rows use 2000 burn-in iterations and 3000 retained draws. Within each model block, bold entries indicate the smaller value for that metric.

values. In practice we recommend selecting the discount factors with the CRPS and using the KL divergence as a complementary calibration check. A simple example is seen below.

```
R> possible.dfs = cbind(0.9,seq(0.85,1,0.05))
R> possible.dfs

     [,1] [,2]
[1,]  0.9 0.85
[2,]  0.9 0.90
[3,]  0.9 0.95
[4,]  0.9 1.00

R> metrics <- matrix(NA_real_, nrow(possible.dfs), 2,
+                    dimnames = list(NULL, c("CRPS","KL")))
R> ref.samp = rnorm(length(sunspot.year))
R> for(i in 1:nrow(possible.dfs)){
+   temp.M2 = exdqlmLDVB(y = sunspot.year, p0 = 0.85, model = model,
+                  df = possible.dfs[i,], dim.df = c(1,8),
+                  sig.init = 2, n.samp = 3000, verbose = FALSE)
+   temp.check = exdqlmDiagnostics(temp.M2, plot = FALSE,
+                              ref = ref.samp)
+   metrics[i,] = c(temp.check$m1.CRPS, temp.check$m1.KL)
+ }
R> # optimal dfs based off CRPS
R> possible.dfs[which.min(metrics[, "CRPS"]),]

[1] 0.90 0.85
```

On this coarse grid, the CRPS and KL criteria agree and both select the seasonal discount factor 0.85, so the choice used above is retained.

The runtime comparisons below were obtained under the default implementation used by the dynamic routines. Exact hardware, software, seed, and dataset details are recorded in the accompanying reproducibility materials.

Table **??** gives a compact dynamic runtime-and-quality comparison under matched 3000-draw summaries. For the simpler DQLM, the VB fit remains much faster than MCMC, whereas the MCMC fit achieves smaller values for all three reported diagnostics. For the more flexible exDQLM, the VB fit is also faster than the corresponding MCMC run, but the comparison is mixed: the VB fit attains the smallest KL divergence in the table, while the MCMC fit

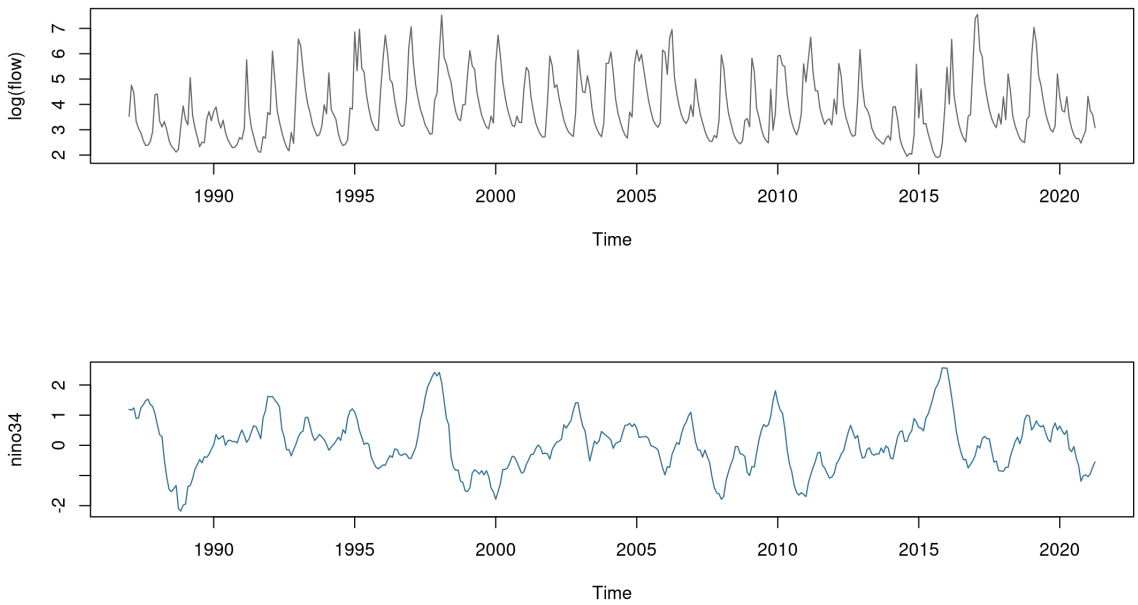

Figure 5: Example 3: Big Tree water flow. Top: average monthly water flow at the Big Tree gauge of the San Lorenzo river in Santa Cruz County, CA, plotted on the log scale after aggregating the daily series to monthly means. Bottom: monthly values of Niño 3.4. Both series span January 1987 through April 2021.

is preferred by both the CRPS and the pplc. Within each inferential method, the exDQLM improves the reported diagnostics relative to the DQLM. These results again reinforce the importance of reporting runtime together with predictive criteria rather than runtime alone: the practical tradeoff depends on both the model and the inferential method being compared.

4.3. Big Tree water flows

For our final example, we consider the average monthly natural water flow (cubic feet per second) at the Big Tree gauge of the San Lorenzo river in Santa Cruz County, CA, obtained by averaging daily USGS observations within each calendar month (**?**). The resulting monthly series is included in the package as `BTflowUSGS`. To keep the covariate structure of the original example, we pair this monthly USGS-derived series with the monthly Niño 3.4 index (**?**). Because the package `nino34` series ends in April 2021, we restrict attention to the overlapping window from January 1987 through April 2021, for a total of 412 observations. With this example we show how to build a state-space structure manually for a dynamic quantile regression, apply our transfer-function exDQLM, examine individual components of the estimated quantiles, and perform quantitative model diagnostics to compare models.

```
R> flow.monthly = window(BTflowUSGS, end = c(2021,4))
R> nino.monthly = window(nino34, start = start(flow.monthly),
+                        end = end(flow.monthly))
R> flow.fit = flow.monthly
```

```
R> nino.input = nino.monthly
```

For this particular dataset, the interest is in low quantiles as a possible indicator of drought, thus we consider the $p_0 = 0.15$ quantile. To investigate climatological influence on this low quantile, we include the monthly Niño 3.4 index as an input covariate in both models. The dynamic regression specification uses Niño 3.4 directly as a time-varying regressor, while the transfer-function model uses the same contemporaneous monthly input and lets the transfer-function structure capture persistence through ζ_t and λ. The fitted comparison below therefore uses the full overlapping sample from January 1987 through April 2021, for a total of 412 observations. To avoid modeling problems caused by the scale and extremes in the Big Tree dataset, we fit the following models to the log monthly flow series shown in Figure **??**.

In addition to the component that captures the relationship with the Niño 3.4 input, we include a first-order polynomial trend component together with a Fourier-form seasonal component with period 12. To mirror the annualized seasonal structure used in the daily redesign, we include harmonics $h = 1$, $h = 2$, and $h = 0.1469118636$, corresponding to annual, semiannual, and approximately 6.8-year cycles on the monthly scale.

```
R> trend.comp = polytrendMod(1, m0 = 3, C0 = 0.1)
R> seas.comp = seasMod(p = 12, h = c(1, 2, 0.1469118636),
+                      C0 = diag(1,6))
R> model = trend.comp + seas.comp
```

Next, we create the dynamic regression component for our first model using the monthly Niño 3.4 series directly. Combining this with the trend and seasonal components gives our complete state-space structure for the dynamic quantile regression model.

```
R> reg.comp = regMod(nino.input, m0 = 0, C0 = matrix(1,1,1))
R> model.w.reg = model + reg.comp
```

To allow all components to evolve gradually over time while keeping the fitted paths smooth, we set every discount factor in both models to 0.97. In particular, we use 0.97 for the trend block, the seasonal block, the monthly Niño 3.4 regression block, and the transfer-function block.

The transfer-function model follows the structure detailed in Section **??**. For this illustration we use `exdqlmTransferLDVB()` to obtain a fast approximate fit on the monthly series, although the corresponding exact posterior workflow is also available through `exdqlmTransferMCMC()`. On the coarse candidate grid used here, we use the KL divergence as a calibration-oriented screening criterion to select the transfer-function decay rate, and then assess the resulting fitted models with the broader diagnostic comparison reported below. This screen selects a decay rate of 0.50.

```
R> possible.lams = c(0.1,0.3,0.5,0.7,0.85)
R> KLs <- vector("numeric")
R> ref.samp = rnorm(length(flow.fit))
R> for(i in 1:length(possible.lams)){
R>   temp.M2 = exdqlmTransferLDVB(y = log(flow.fit), p0 = 0.15, model = model,
+                                df = c(0.97,0.97), dim.df = c(1,6),
+                                X = nino.input, tf.df = c(0.97),
+                                lam = possible.lams[i],
+                                tf.m0 = c(0,0),
```

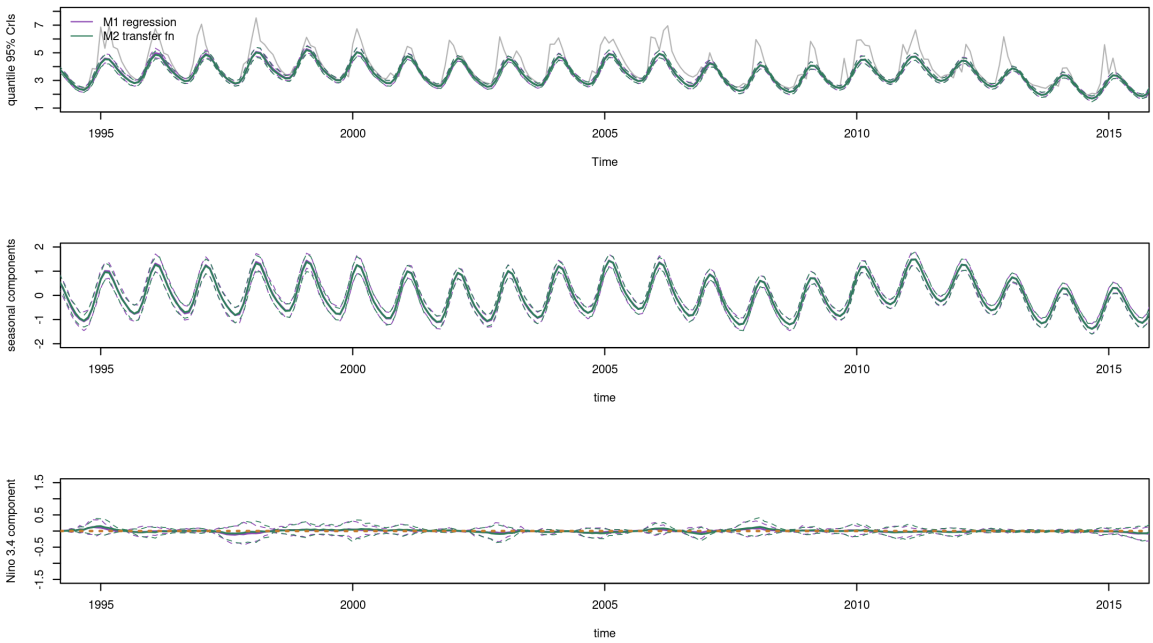

Figure 6: Example 3: Big Tree water flow. Top: MAP estimates and 95% CrIs of the estimated quantiles from the models with the dynamic regression component (purple) and transfer-function component (green), plotted with the data (grey) from 1995 to 2015. Middle: MAP estimates and 95% CrIs of the combined seasonal components implied by the annual, semiannual, and low-frequency harmonics. Bottom: MAP estimates and 95% CrIs of the components that model the quantile association with Niño 3.4. For M_1, the purple curve is the regression contribution of the monthly Niño 3.4 covariate. For M_2, the green curve is the transfer-function contribution. The horizontal orange dotted line is at zero for reference.

```
+                                tf.C0 = diag(c(0.1,0.005),2),
+                                sig.init = 0.1, gam.init = - 0.1,
+                                tol = 0.05, verbose = FALSE)
R>   temp.check = exdqlmDiagnostics(temp.M2, plot = FALSE,
+                              ref = ref.samp)
R>   KLs = c(KLs,temp.check$m1.KL)
R> }
R> # optimal lambda based off KL divergence
R> possible.lams[which.min(KLs)]

[1] 0.5
```

With these discount factors and optimal λ, we estimate the dynamic quantile regression model (`M1`) and transfer function model (`M2`).

```
R> M1 = exdqlmLDVB(y = log(flow.fit), p0 = 0.15, model = model.w.reg,
+                 df = c(0.97,0.97,0.97), dim.df = c(1,6,1),
+                 sig.init = 0.1, gam.init = - 0.1,
+                 tol = 0.05, verbose = FALSE)
```

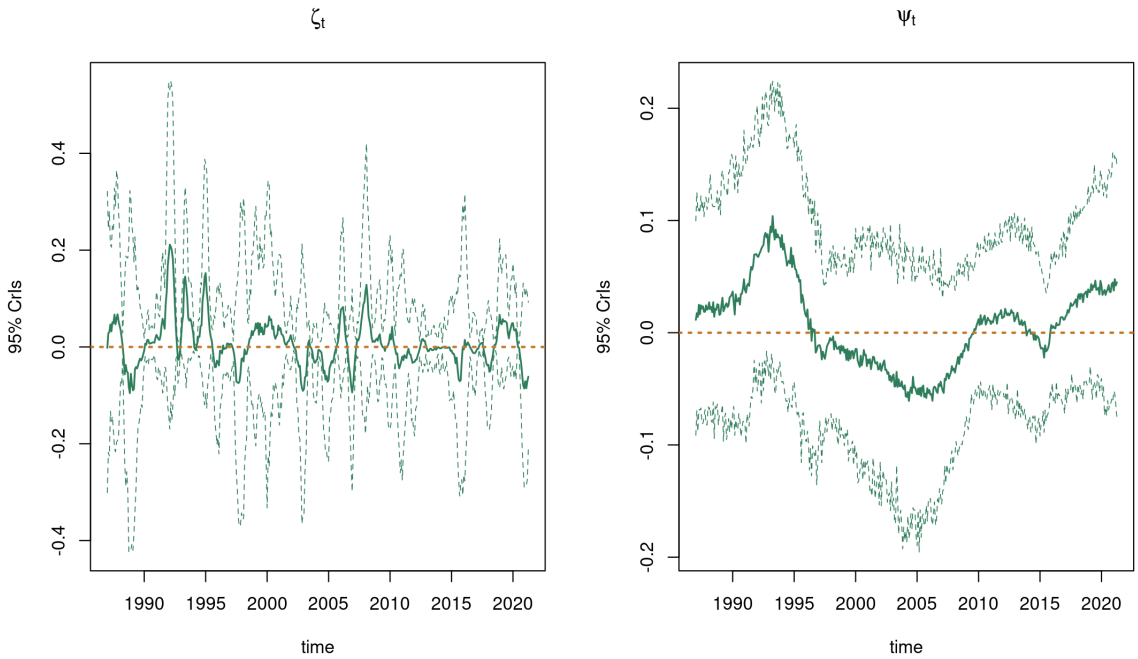

Figure 7: Example 3: Big Tree water flow. MAP estimates and 95% CrIs of the transfer-function states ζ_t and ψ_t. Horizontal dotted orange lines are at zero for reference.

```
R> M2 = exdqlmTransferLDVB(y = log(flow.fit), p0 = 0.15, model = model,
+                          df = c(0.97,0.97), dim.df = c(1,6),
+                          X = nino.input, tf.df = c(0.97), lam = 0.50,
+                          tf.m0 = c(0,0),
+                          tf.C0 = diag(c(0.1,0.005),2),
+                          sig.init = 0.1, gam.init = - 0.1,
+                          tol = 0.05, verbose = FALSE)
```

To visualize the differences between the two resulting dynamic quantiles we use the function `exdqlmPlot`, the results of which are found in Figure **??**. For clarity we limit the figure to the period 1995 to 2015.

```
R> plot(log(flow.fit), col = "grey", ylim = c(1,8), xlim=c(1995,2015),
+       ylab = "quantile 95% CrIs")
R> exdqlmPlot(M1, add = TRUE)
R> exdqlmPlot(M2, add = TRUE, col = "forest green")
```

The resulting quantiles exhibit visible differences. To identify from where these differences arise, we can use the function `compPlot` to examine the model components more closely. We examine the combined seasonal block first. Because the seasonal structure now contains three harmonics, the seasonal contribution is formed by the second through seventh elements of the state vector. The result is seen in Figure **??**, again limited to the years 1995 to 2015.

```
R> plot(NA, ylim = c(-2,2), xlim = c(1995,2015),
+         ylab = "seasonal components")
```

```
R> compPlot(M1, index = 2:7, add = TRUE)
R> compPlot(M2, index = 2:7, add = TRUE, col = "forest green")
```

Although the seasonal components of the two models exhibit differences in amplitude, more substantial differences appear in the Niño 3.4-related components. Again, we use the function `compPlot` to visualize these components seen in Figure **??** from 1995 to 2015. For the regression model, the plotted contribution is simply the dynamic regression contribution of the monthly Niño 3.4 covariate.

```
R> plot(NA, ylim = c(-2,2), xlim = c(1995,2015),
+          ylab = "Nino 3.4 components")
R> compPlot(M1, index = 8, add = TRUE)
R> compPlot(M2, index = 8:9, add = TRUE, col = "forest green")
R> abline(h = 0, col = "orange", lty = 3, lwd = 2)
```

For the transfer-function model, `M2`, it is also of interest to look at the estimates of parameters ζ_t and ψ_t directly. Because these parameters augment the state-space model as discussed in Section **??**, the function `compPlot` can also be used to look at them directly via the parameter `just.theta = TRUE`. This setting returns the dynamic posterior estimates of a single index of the state vector $\boldsymbol{\theta}_t$ (or $\tilde{\boldsymbol{\theta}}_t$ in the case of the transfer-function model). Here the indices 8 and 9 correspond to ζ_t and ψ_t, respectively. Results can be found in Figure **??**.

```
R> compPlot(M2, index = 8, col= "forest green",
+             add = FALSE, just.theta = TRUE)
R> abline(h = 0, col = "orange", lty = 3, lwd = 2)
R> title(expression(zeta[t]))
R> compPlot(M2, index = 9, col = "forest green",
+             add = FALSE, just.theta = TRUE)
R> abline(h = 0, col = "orange", lty = 3, lwd = 2)
R> title(expression(psi[t]))
```

To further understand these transfer-function parameters, we examine the median number of time steps until the effect of Niño 3.4 on the 0.15 quantile is less than or equal to 1e−3, as discussed in Section **??**. Here the effect is $|\psi_{t-1}X_t|$, where X_t denotes the contemporaneous monthly Niño 3.4 input. A median value of about 3.76 suggests that the transfer effect lasts on the order of four months.

```
R> M2$median.kt
```

```
[1] 3.758485
```

Because the Niño 3.4 series is available through April 2021, we use the final 18 months of the overlapping sample to illustrate the forecasting workflow; the results are shown in Figure **??**.

```
R> exdqlmForecast(start.t = length(flow.monthly)-18, k = 18, m1 = M1)
R> exdqlmForecast(start.t = length(flow.monthly)-18, k = 18, m1 = M2,
+                   add = TRUE, cols = c("forest green","green"))
```

Finally, we use the function `exdqlmDiagnostics()` to compute the KL divergence, CRPS, and pplc discussed in Section **??**. Table **??** summarizes the resulting diagnostic comparison.

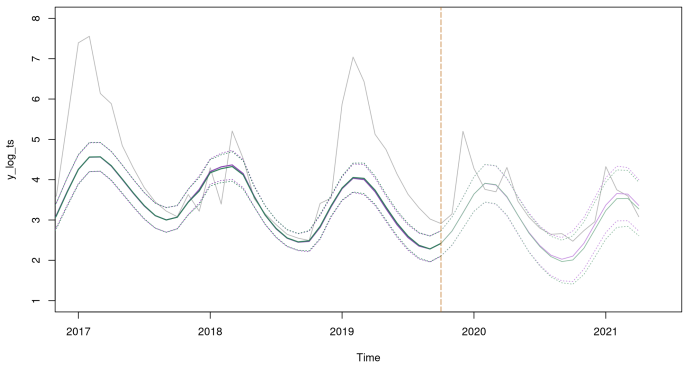

Figure 8: Example 3: Big Tree water flow. 18-step-ahead quantile forecast beginning October 2019 and extending through April 2021, overlaid on the log Big Tree water flow time series. The vertical dot-dashed line marks the start of the forecast period. Solid lines indicate mean estimates and dotted lines 95% CrI, with the filtered estimates shown before the forecast origin and the forecast estimates beyond it. Estimates from M_1, the quantile regression model, are seen in purple/pink. Estimates from M_2, the transfer function model, are seen in green.

Criterion	`M1` regression	`M2` transfer function
CRPS	**0.350**	0.357
pplc	132.10	**122.90**
KL	**0.0562**	0.0565

Table 8: Example 3: Big Tree water flow. Diagnostic comparison between the dynamic regression model (`M1`) and the transfer-function model (`M2`). Lower values are preferred for all three criteria. Here CRPS assesses the full predictive distribution, pplc assesses target-quantile posterior predictive loss, and KL summarizes one-step-ahead calibration discrepancy. Bold entries indicate the smaller value in each row.

```
R> checks.out = exdqlmDiagnostics(M1, M2,
+                                 cols = c("purple","orange"), plot = FALSE)
```

With this simpler monthly Niño 3.4 specification, the resulting diagnostics are again not unanimous. The direct regression model attains the smaller KL divergence and CRPS, whereas the transfer-function model achieves the smaller pplc. The differences are fairly modest, so we view the two specifications as offering closely competing summaries of the same low-quantile dynamics: the direct regression is slightly favored by the calibration-oriented and full-distribution criteria, while the transfer-function model is favored by the target-quantile loss criterion.

4.4. Static exAL regression on a simulated heteroskedastic example

To illustrate the static routines under sparse shrinkage, we consider a correlated Gaussian regression benchmark for which the target conditional quantile is known exactly. Let

$$\mathbf{x}_i \sim N_8(\mathbf{0}, \Sigma_x), \qquad (\Sigma_x)_{jk} = 0.5^{|j-k|},$$

with sparse signal

$$\boldsymbol{\beta}^\star = (3,\ 1.5,\ 0,\ 0,\ 2,\ 0,\ 0,\ 0)^\top.$$

We use the following benchmark data-generating mechanism. For each target quantile level p_0, we generate

$$Y_i^{(p_0)} = \mathbf{x}_i^\top \boldsymbol{\beta}^\star + \sigma_\varepsilon\{Z_i - \Phi^{-1}(p_0)\}, \qquad Z_i \sim N(0,1), \qquad \sigma_\varepsilon = 1.5.$$

Then

$$Q_Y(p_0 \mid \mathbf{x}_i) = \mathbf{x}_i^\top \boldsymbol{\beta}^\star,$$

so the true p_0-quantile is the same sparse linear signal at each fitted quantile level. We use $n = 160$ training observations and an independent holdout sample of size 800, and fit separate static exAL models at $p_0 \in \{0.05, 0.25, 0.50\}$. Both fits use the regularized horseshoe (RHS) prior of **?**, implemented in the package through `beta_prior = "rhs_ns"`, with $\tau_0 = 0.15$, fixed slab $\zeta^2 = 9$, and an unshrunk intercept. We fit a VB approximation with `exalStaticLDVB()` using `n.samp = 3000`, and an MCMC fit with `exalStaticMCMC()` using `n.burn = 2000` and `n.mcmc = 3000`; the MCMC routine uses its default `slice` kernel and is warm-started from the VB fit. The lower-tail VB fit at $p_0 = 0.05$ converges more slowly, so we allow a larger VB iteration budget there. For reproducibility, we fix the dataset seed at `20260714`. Setting `dqlm.ind = TRUE` would recover the AL special case discussed in Section **??**, but here we focus on the general static exAL model under sparse RHS shrinkage.

```
R> Sigma.x = 0.5 ^ as.matrix(dist(1:8))
R> beta.true = c(3, 1.5, 0, 0, 2, 0, 0, 0)
R> rhs.ctrl = list(tau0 = 0.15, a_zeta = 2, b_zeta = 9,
                   zeta2_fixed = 9, shrink_intercept = FALSE)
R> set.seed(20260714)
R> sim.y = function(X.raw, p0) {
      drop(X.raw %*% beta.true) + 1.5 * (rnorm(nrow(X.raw)) - qnorm(p0))
    }
R> X.raw = MASS::mvrnorm(160, mu = rep(0, 8), Sigma = Sigma.x)
R> X = cbind(1, X.raw)
R> p.grid = c(0.05, 0.25, 0.50)
R> y.train = lapply(p.grid, function(p0) sim.y(X.raw, p0))
R> M.ldvb = lapply(p.grid, function(p0) exalStaticLDVB(
+                    y = y.train[[which(p.grid == p0)]], X = X, p0 = p0,
+                    beta_prior = "rhs_ns",
+                    beta_prior_controls = rhs.ctrl,
+                    max_iter = ifelse(p0 == 0.05, 420, 260),
+                    n.samp = 3000,
+                    tol = 1e-4, verbose = FALSE))
R> M.mcmc = lapply(p.grid, function(p0) exalStaticMCMC(
+                    y = y.train[[which(p.grid == p0)]], X = X, p0 = p0,
+                    beta_prior = "rhs_ns",
+                    beta_prior_controls = rhs.ctrl,
+                    n.burn = 2000, n.mcmc = 3000,
+                    init.from.vb = TRUE,
+                    verbose = FALSE))
```

For static fits, the helper function `exalStaticDiagnostics()` summarizes fitted quantiles on a common design matrix and, when a reference quantile is supplied, reports truth-recovery metrics directly. For example, on an independent holdout design:

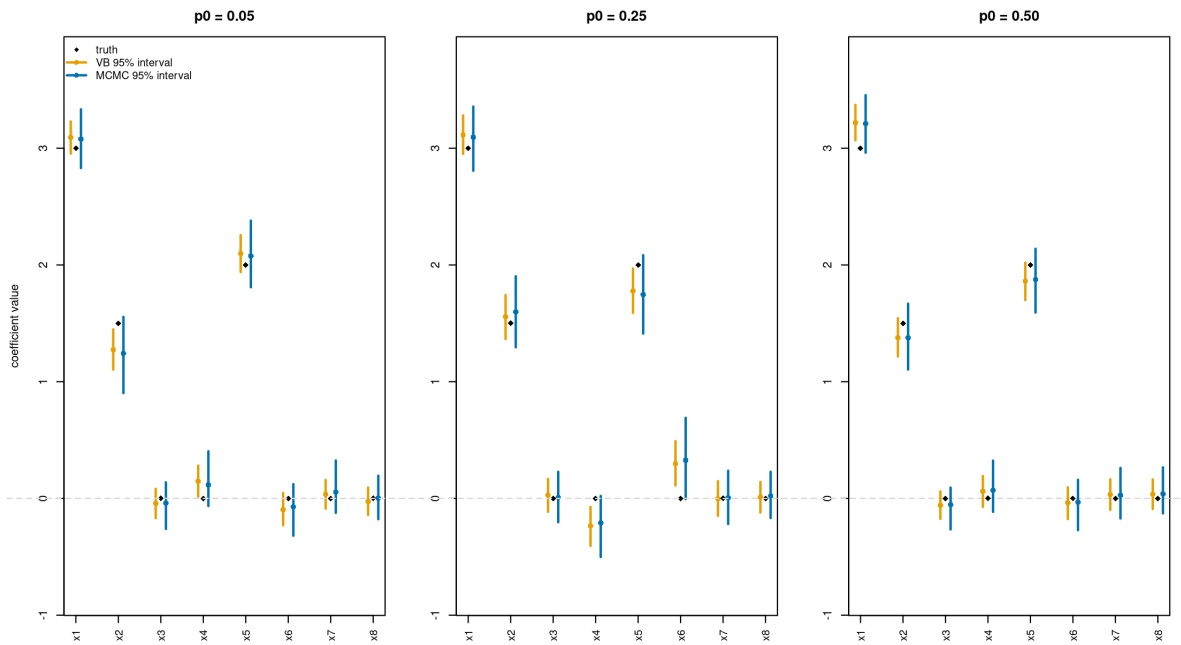

Figure 9: Example 4: sparse static exAL regression under the correlated Gaussian benchmark with the RHS prior. Black diamonds denote the true non-intercept coefficients, orange intervals denote VB 95% posterior intervals, and blue intervals denote MCMC 95% posterior intervals. Panels correspond to p0 = 0.05, 0.25, and 0.50.

```
R> X.hold.raw = MASS::mvrnorm(800, mu = rep(0, 8), Sigma = Sigma.x)
R> X.hold = cbind(1, X.hold.raw)
R> ref.hold = drop(X.hold %*% c(0, beta.true))
R> y.hold = sim.y(X.hold.raw, 0.25)
R> exalStaticDiagnostics(M.ldvb[[2]], M.mcmc[[2]],
+                   X = X.hold, y = y.hold, ref = ref.hold,
+                   plot = FALSE)
```

Figure **??** shows that both methods recover the active coefficients and shrink the null coefficients toward zero across the lower tail, lower-middle quantile, and median. The lower-tail fit at $p_0 = 0.05$ is the most demanding case, yet both methods still recover the sparse signal cleanly. Across all three quantile levels, the visual differences between VB and MCMC are modest, with the most noticeable separation appearing in the lower tail.

Table **??** reinforces the visual comparison. The two methods remain competitive across all three quantile levels, but the relative advantage depends on the metric. At $p_0 = 0.05$, the VB fit is slightly better on the active coefficients, whereas MCMC is better on the null coefficients and the holdout quantile RMSE. At $p_0 = 0.25$, the VB fit is slightly better on the active-signal and holdout metrics, while the null-coefficient summaries are effectively tied. At $p_0 = 0.50$, MCMC again holds a small advantage on the active-signal and holdout metrics, and the null-coefficient summaries are nearly indistinguishable. The VB fit is nevertheless much faster in all three cases, with runtime reductions of roughly 7-fold at $p_0 = 0.05$, 11-fold at $p_0 = 0.25$, and 14-fold at $p_0 = 0.50$. In this benchmark, the VB approximation offers meaningful

p_0	method	runtime (s)	active RMSE	null MAE	holdout quantile RMSE
0.05	VB	**19.11**	**0.153**	0.069	0.471
	MCMC	136.99	0.162	**0.057**	**0.419**
0.25	VB	**11.45**	**0.148**	0.115	**0.401**
	MCMC	125.99	0.166	**0.114**	0.417
0.50	VB	**9.37**	0.166	**0.045**	0.247
	MCMC	127.27	**0.159**	**0.045**	**0.238**

Table 9: Example 4: runtime and sparse-signal recovery metrics for the static exAL fits under the RHS prior. The VB rows use `n.samp = 3000`, and the MCMC rows use 2000 burn-in iterations and 3000 retained draws. The active RMSE is computed over the three nonzero coefficients, the null MAE is computed over the five zero coefficients, and the holdout quantile RMSE is computed against the true target quantile on an independent holdout sample of size 800. Within each quantile block, bold entries indicate the smaller value for that metric; ties are bolded in both rows.

computational savings as a first-pass analysis while still recovering the sparse signal well, but this speed comes with the usual variational tradeoff that posterior uncertainty quantification may be less faithful than the exact MCMC baseline. Accordingly, the static RHS specification supports both roles cleanly: VB for rapid screening and model exploration, and static MCMC as the reference method when higher-fidelity posterior uncertainty quantification is required.

5. Conclusion

In this paper we presented the R package **exdqlm** for Bayesian estimation and analysis of flexible dynamic quantile linear models. The package provides exact posterior inference via MCMC and fast approximate posterior inference via VB for dynamic state-space quantile models, while also supporting transfer-function exDQLMs, static exAL quantile regression with shrinkage priors, forecasting, model diagnostics, and post hoc posterior-predictive synthesis across separately fitted quantiles. Together, these functions provide a practical framework for dynamic and static quantile modeling beyond the AL special case while preserving a computationally convenient latent-variable representation.

The examples illustrate the main workflows supported by the package. The dynamic examples show how **exdqlm** can be used to construct state-space models, estimate dynamic quantiles, compare the exDQLM with its AL special case, examine transfer function effects, and evaluate model fit through diagnostics and forecasting. The static example shows how the same framework can be used for sparse quantile regression with both VB and MCMC, with VB serving as a fast first-pass approximation and MCMC providing the reference analysis when more faithful posterior uncertainty quantification is required. Viewed together, these examples highlight the two software contributions that most clearly distinguish the package within the current ecosystem: Bayesian dynamic state-space quantile modeling and a unified workflow that extends from static regularized exAL regression to post hoc predictive synthesis.

References

Barata R, Prado R, Sansó B (2022). "Fast inference for time-varying quantiles via flexible dynamic models with application to the characterization of atmospheric rivers." *The annals of applied statistics*.

Barlow RE, Bartholomew DJ, Bremner JM, Brunk HD (1972). *Statistical Inference Under Order Restrictions: The Theory and Application of Isotonic Regression*. John Wiley & Sons, New York.

Beal MJ (2003). *Variational algorithms for approximate Bayesian inference*. Ph.D. thesis, UCL (University College London).

Benoit D, Al-Hamzawi R, Yu K, Van den Poel D (2023). ***bayesQR**: Bayesian Quantile Regression*. R package version 2.4, URL https://CRAN.R-project.org/package=bayesQR.

Bürkner PC (2026). ***brms**: Special Family Functions for **brms** Models*. Package documentation, accessed 2026-04-19, URL https://paulbuerkner.com/brms/reference/brmsfamily.html.

Carter CK, Kohn R (1994). "On Gibbs sampling for state space models." *Biometrika*, **81**(3), 541–553.

Chernozhukov V, Fernández-Val I, Galichon A (2010). "Quantile and Probability Curves Without Crossing." *Econometrica*, **78**(3), 1093–1125. doi:10.3982/ECTA7880.

Fan K, Wu C, Ren J, Li X, Zhou F (2026). ***pqrBayes**: Bayesian Penalized Quantile Regression*. R package version 1.2.1, URL https://CRAN.R-project.org/package=pqrBayes.

Fasiolo M, Griffiths B (2025). ***qgam**: Smooth Additive Quantile Regression Models*. R package version 2.0.0, URL https://CRAN.R-project.org/package=qgam.

Fogarty B (2020). "**QuantReg.jl**: Quickstart Guide." Documentation page, accessed 2026-04-19, URL https://fogarty-ben.github.io/QuantReg.jl/v0.1/quickstart/.

Fox K (2012). "Solar Minimum; Solar Maximum." URL https://www.nasa.gov/mission_pages/sunearth/news/solarmin-max.html.

Frühwirth-Schnatter S (1994). "Data augmentation and dynamic linear models." *Journal of time series analysis*, **15**(2), 183–202.

Frumento P (2025). ***qrcm**: Quantile Regression Coefficients Modeling*. R package version 3.2, URL https://CRAN.R-project.org/package=qrcm.

Galarza CE, Benites L, Bourguignon M, Lachos VH (2024). ***lqr**: Robust Linear Quantile Regression*. R package version 5.2, URL https://CRAN.R-project.org/package=lqr.

Gelfand AE, Ghosh SK (1998). "Model choice: a minimum posterior predictive loss approach." *Biometrika*, **85**(1), 1–11.

Gneiting T, Raftery AE (2007). "Strictly Proper Scoring Rules, Prediction, and Estimation." *Journal of the American Statistical Association*, **102**(477), 359–378. doi:10.1198/016214506000001437.

Gonçalves K, Migon HS, Bastos LS (2017). "Dynamic quantile linear model: a Bayesian approach." *arXiv preprint arXiv:1711.00162*.

He X, Pan X, Tan KM, Zhou WX (2023). ***conquer**: Convolution-Type Smoothed Quantile Regression*. R package version 1.3.3, URL https://CRAN.R-project.org/package=conquer.

Huerta G, Jiang W, Tanner MA (2003). "Time series modeling via hierarchical mixtures." *Statistica Sinica*, pp. 1097–1118.

Koenker R (2005). *Quantile regression*. Cambridge University Press, New York.

Koenker R (2025). ***quantreg**: Quantile Regression*. R package version 6.1, URL https://CRAN.R-project.org/package=quantreg.

Kottas A, Krnjajić M (2009). "Bayesian semiparametric modelling in quantile regression." *Scandinavian Journal of Statistics*, **36**(2), 297–319.

Kozumi H, Kobayashi G (2011). "Gibbs sampling methods for Bayesian quantile regression." *Journal of statistical computation and simulation*, **81**(11), 1565–1578.

Kullback S, Leibler RA (1951). "On information and sufficiency." *The annals of mathematical statistics*, **22**(1), 79–86.

MathWorks (2026). "Working with Quantile Regression Models." Documentation page, accessed 2026-04-19, URL https://www.mathworks.com/help/stats/working-with-quantile-regression-models.html.

Muggeo VMR (2025). ***quantregGrowth**: Non-Crossing Additive Regression Quantiles and Non-Parametric Growth Charts*. R package version 1.7-2, URL https://CRAN.R-project.org/package=quantregGrowth.

Nishimura A, Suchard MA (2023). "Shrinkage with shrunken shoulders: Gibbs sampling shrinkage model posteriors with guaranteed convergence rates." *Bayesian Analysis*, **18**(2), 367–390. doi:10.1214/22-BA1308.

Ostwald D, Kirilina E, Starke L, Blankenburg F (2014). "A tutorial on variational Bayes for latent linear stochastic time-series models." *Journal of Mathematical Psychology*, **60**, 1–19.

Ou L, Hunter M, Chow SM, Ji L, Chen M, Hung HJ, Lee J, Li Y, Park J (2021). ***dynr**: Dynamic Models with Regime-Switching*. R package version 0.1.16-2, URL https://CRAN.R-project.org/package=dynr.

Petris G, Gilks W (2018). ***dlm**: Bayesian and Likelihood Analysis of Dynamic Linear Models*. R package version 1.1-5, URL https://CRAN.R-project.org/package=dlm.

Prado R, Molina F, Huerta G (2006). "Multivariate time series modeling and classification via hierarchical VAR mixtures." *Computational Statistics & Data Analysis*, **51**(3), 1445–1462.

Rayner N, Parker DE, Horton E, Folland CK, Alexander LV, Rowell D, Kent EC, Kaplan A (2003). "Global analyses of sea surface temperature, sea ice, and night marine air temperature since the late nineteenth century." *Journal of Geophysical Research: Atmospheres*, **108**(D14).

Reich BJ, Bondell HD, Wang HJ (2009). "Flexible Bayesian quantile regression for independent and clustered data." *Biostatistics*, **11**(2), 337–352.

Rosenblatt M (1952). "Remarks on a multivariate transformation." *The annals of mathematical statistics*, **23**(3), 470–472.

Sherwood B, Li S, Maidman A (2026). ***rqPen**: Penalized Quantile Regression*. R package version 4.2, URL https://CRAN.R-project.org/package=rqPen.

statsmodels developers (2025). "**statsmodels**: `QuantReg` Documentation." Version 0.14.6 documentation, accessed 2026-04-19, URL https://www.statsmodels.org/stable/generated/statsmodels.regression.quantile_regression.QuantReg.html.

Taddy MA, Kottas A (2010). "A Bayesian nonparametric approach to inference for quantile regression." *Journal of Business & Economic Statistics*, **28**(3), 357–369.

Tokdar S, Cunningham E (2025). ***qrjoint**: Joint Estimation in Linear Quantile Regression*. R package version 2.0-11, URL https://CRAN.R-project.org/package=qrjoint.

US Geological Survey (2016). "National Water Information System data available on the World Wide Web (USGS Water Data for the Nation)." URL http://waterdata.usgs.gov/nwis/.

Wang C, Blei DM (2013). "Variational Inference in Non-Conjugate Models." *arXiv preprint arXiv:1209.4360*. URL https://arxiv.org/abs/1209.4360.

West M, Harrison PJ, Migon HS (1985). "Dynamic generalized linear models and Bayesian forecasting." *Journal of the American Statistical Association*, **80**(389), 73–83.

Yan Y, Zheng X, Kottas A (2025). "A New Family of Error Distributions for Bayesian Quantile Regression." *Bayesian Analysis*. doi:10.1214/25-BA1507.

Yu K, Moyeed RA (2001). "Bayesian quantile regression." *Statistics & Probability Letters*, **54**(4), 437–447.

Affiliation:

Antonio Aguirre and Raquel Barata
Department of Statistics
University of California Santa Cruz
Santa Cruz, CA 95064
E-mail: jaguir26@ucsc.edu, raquel.a.barata@gmail.com

Journal of Statistical Software http://www.jstatsoft.org/
published by the Foundation for Open Access Statistics http://www.foastat.org/
MMMMMM YYYY, Volume VV, Issue II *Submitted:* yyyy-mm-dd
doi:10.18637/jss.v000.i00 *Accepted:* yyyy-mm-dd